

# Homework3\_CSCE633

November 12, 2021

## 1 Homework\_3\_432001358\_CSCE\_633\_600

### 1.1 November 12, 2021

### 1.2 CSCE 633 600 (Machine Learning) Homework 3

#### 1.2.1 Name: Rohan Chaudhury

#### 1.2.2 UIN: 432001358

#### 1.2.3 Question 1: Machine learning for facial emotion recognition

In this problem, we will process face images coming from the Facial Expression Recognition Challenge (presented in the International Conference of Machine Learning in 2013). The data is uploaded under Homework3 folder in the shared Google Drive. You are given three sets of data: training set (i.e., Q1 Train Data.csv), testing set (i.e., Q1 Test Data.csv), and validation set (i.e., Q1 Validation Data.csv). The data consists of 48X48 pixel grayscale images of faces. The faces have been automatically registered so that the face is more or less centered and occupies about the same amount of space in each image. The task is to categorize each face based on the emotion shown in the facial expression in seven categories. More information on the data can also be found in this link. All three files contain two columns: 1. The column labeled as “emotion” contains the emotion class with numeric code ranging from 0 to 6 (0=Angry, 1=Disgust, 2=Fear, 3=Happy, 4=Sad, 5=Surprise, 6=Neutral). 2. The column labeled as “pixels” contains the 2304 (i.e., 48 X 48) space-separated pixel values of the image in row-wise order, i.e., the first 48 numbers correspond to the first row of the image, the next 48 numbers to the second row of the image, etc.

```
[2]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import tensorflow as tf

train_data = pd.read_csv('Q1_Train_Data.csv')
test_data = pd.read_csv('Q1_Test_Data.csv')
validation_data = pd.read_csv('Q1_Validation_Data.csv')

print (" ")
print("Shape of Train Data: {}".format(train_data.shape))
print (" ")
```

```
print("Shape of Test Data: {}".format(test_data.shape))
print (" ")
print("Shape of Validation Data: {}".format(validation_data.shape))
print (" ")
print(train_data.head)
```

Shape of Train Data: (28709, 2)

Shape of Test Data: (3589, 2)

Shape of Validation Data: (3589, 2)

```
<bound method NDFrame.head of          emotion
pixels
0          0   70 80 82 72 58 58 60 63 54 58 60 48 89 115 121...
1          0  151 150 147 155 148 133 111 140 170 174 182 15...
2          2  231 212 156 164 174 138 161 173 182 200 106 38...
3          4   24 32 36 30 32 23 19 20 30 41 21 22 32 34 21 1...
4          6   4 0 0 0 0 0 0 0 0 0 0 0 3 15 23 28 48 50 58 84...
...      ...
28704      2   84 85 85 85 85 85 85 85 86 86 86 87 86 86 91 9...
28705      0  114 112 113 113 111 111 112 113 115 113 114 11...
28706      4   74 81 87 89 95 100 98 93 105 120 127 133 146 1...
28707      0  222 227 203 90 86 90 84 77 94 87 99 119 134 14...
28708      4  195 199 205 206 205 203 206 209 208 210 212 21...
```

[28709 rows x 2 columns]>

**1.2.4 (a) (1 points) Visualization:** Randomly select and visualize 1-2 images per emotion. Note: You can find a useful link on image pre-processing here: [https://www.tensorflow.org/api\\_docs/python/tf/image/per\\_image\\_standardization](https://www.tensorflow.org/api_docs/python/tf/image/per_image_standardization)

```
[3]: def string_to_int(st):
      ar= list((int(x)/255 - 0.5) for x in st.split(' '))
      return ar

      def get_pixels_int(df):
          df['pixels_int']=df.apply(lambda st: np.
          ↳array(string_to_int(st['pixels'])),axis=1)

      get_pixels_int(train_data)
      get_pixels_int(test_data)
      get_pixels_int(validation_data)

      print(train_data.head)
```

```

<bound method NDFrame.head of          emotion ...
pixels_int
0          0 ... [-0.22549019607843135, -0.18627450980392157, -...
1          0 ... [0.09215686274509804, 0.08823529411764708, 0.0...
2          2 ... [0.40588235294117647, 0.3313725490196079, 0.11...
3          4 ... [-0.40588235294117647, -0.37450980392156863, -...
4          6 ... [-0.4843137254901961, -0.5, -0.5, -0.5, -...
...      ...
28704      2 ... [-0.17058823529411765, -0.16666666666666669, -...
28705      0 ... [-0.052941176470588214, -0.0607843137254902, -...
28706      4 ... [-0.20980392156862743, -0.1823529411764706, -0...
28707      0 ... [0.37058823529411766, 0.3901960784313725, 0.29...
28708      4 ... [0.2647058823529411, 0.2803921568627451, 0.303...

```

[28709 rows x 3 columns]>

```

[7]: # plt.rcParams["figure.figsize"] = (2,300)

plt.figure(figsize=(400,400))

emotion_category={0:'Angry', 1:'Disgust', 2:'Fear', 3:'Happy', 4:'Sad', 5:
    ↳'Surprise', 6:'Neutral'}


emotions= train_data['emotion'].unique().tolist()

fig, ax=plt.subplots(len(emotions),2,figsize=(15,15))
fig.tight_layout(pad=3.0)
# print (emotions)
plot_num=0
for emotion in emotions:
    imgs=train_data[train_data["emotion"]==emotion]
    samples=imgs.sample(n=2)
    for i in range(samples.shape[0]):
        img=samples.iloc[i]
        ax[int(plot_num/2),plot_num%2].set_title("Emotion label: {}".
    ↳format(emotion_category[emotion]))
        ax[int(plot_num/2),plot_num%2].imshow(np.array(img['pixels_int']).
    ↳reshape(48,48),cmap=plt.get_cmap('gray'))
        plot_num+=1


```

<Figure size 28800x28800 with 0 Axes>


Emotion label: Angry




Emotion label: Angry




Emotion label: Fear




Emotion label: Fear




Emotion label: Sad




Emotion label: Sad




Emotion label: Neutral




Emotion label: Neutral




Emotion label: Happy




Emotion label: Happy




Emotion label: Surprise




Emotion label: Surprise



Emotion label: Disgust



Emotion label: Disgust



1.2.5 (b) (1 points) Data exploration: Count the number of samples per emotion in the training data.

```
[8]: samples_count=train_data.groupby(['emotion']).count()
      samples_count=samples_count[['pixels']].rename({'pixels': 'size'}, axis=1)
      samples_count
```

```
[8]:          size
emotion
0          3995
1           436
2          4097
3          7215
4          4830
5          3171
6          4965
```

1.2.6 (c) (4 points) Image classification with FNNs: In this part, you will use a feed-forward neural network (FNN) (also called multilayer perceptron”) to perform the emotion classification task. The input of the FNN comprises of all the pixels of the image.

```
[9]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense, Dropout
      from tensorflow.keras.utils import to_categorical
      from tensorflow.keras import regularizers

      train_emotions=np.array(list(train_data['emotion']))
      train_pixels=np.array(list(train_data['pixels_int']))

      import time

      class TimeHistory(tf.keras.callbacks.Callback):
          def on_train_begin(self, logs={}):
              self.times = []
              self.epoch_time_start = time.time()
          # def on_epoch_begin(self, epoch, logs={}):
          #     self.epoch_time_start = time.time()

          def on_train_end(self, epoch, logs={}):
              self.times.append(time.time() - self.epoch_time_start)
```

```

model1 = Sequential([
    Dense(784*2, activation='relu', input_shape=(48*48,)),
    ↪name="first_hidden_layer"),
    Dense(784, activation='relu', name="second_hidden_layer"),
    Dense(784//2, activation='relu', name="third_hidden_layer"),
    Dense(784//4, activation='relu', name="fourth_hidden_layer"),
    Dense(len(emotions), activation='softmax'),
])

model2 = Sequential([
    Dense(784*2, activation='elu', input_shape=(48*48,)),
    ↪name="first_hidden_layer"),
    Dense(784, activation='elu', name="second_hidden_layer"),
    Dense(784//2, activation='elu', name="third_hidden_layer"),
    Dropout(0.25),
    Dense(len(emotions), activation='softmax'),
])

model3 = Sequential([
    Dense(2000, activation='elu', input_shape=(48*48,)),
    ↪name="first_hidden_layer"),
    Dense(1000, activation='elu', name="second_hidden_layer"),
    Dense(500, activation='elu', name="third_hidden_layer"),
    Dropout(0.25),
    Dense(len(emotions), activation='softmax'),
])

model4 = Sequential([
    Dense(2000, activation='elu', input_shape=(48*48,)),
    ↪name="first_hidden_layer"),
    Dense(1000, activation='elu', name="second_hidden_layer"),
    Dense(500, activation='elu', name="third_hidden_layer"),
    Dense(250, activation='elu', name="fourth_hidden_layer"),
    Dense(784, activation='relu', name="fifth_hidden_layer"),
    Dropout(0.25),
    Dense(len(emotions), activation='softmax'),
])

model5 = Sequential([
    Dense(2000, activation='elu', input_shape=(48*48,)),
    ↪name="first_hidden_layer", kernel_regularizer=regularizers.l2(0.0001)),
    Dense(1000, activation='elu', name="second_hidden_layer",
    ↪kernel_regularizer=regularizers.l2(0.0001)),

```

```

    Dense(500, activation='elu', name="third_hidden_layer",
    ↪kernel_regularizer=regularizers.l2(0.0001)),
    Dense(250, activation='elu', name="fourth_hidden_layer",
    ↪kernel_regularizer=regularizers.l2(0.0001)),
    Dropout(0.25),
    Dense(len(emotions), activation='softmax'),
])

```

```

[10]: time_callback = TimeHistory()

model1.compile(optimizer='adam',
    ↪loss='categorical_crossentropy',metrics=['accuracy'])
model2.compile(optimizer='adam',
    ↪loss='categorical_crossentropy',metrics=['accuracy'])
model3.compile(optimizer='adam',
    ↪loss='categorical_crossentropy',metrics=['accuracy'])
model4.compile(optimizer='adam',
    ↪loss='categorical_crossentropy',metrics=['accuracy'])
model5.compile(optimizer='adam',
    ↪loss='categorical_crossentropy',metrics=['accuracy'])

print("Train image shape: ", train_pixels.shape)
print(train_emotions.shape)

time_to_train=[]
flatten_train_images = train_pixels

history1=model1.fit(np.array(flatten_train_images),
    ↪to_categorical(train_emotions), epochs=20, batch_size=256,callbacks =
    ↪[time_callback])
time_to_train.append(time_callback.times)

history2=model2.fit(np.array(flatten_train_images),
    ↪to_categorical(train_emotions), epochs=20, batch_size=256,callbacks =
    ↪[time_callback])
time_to_train.append(time_callback.times)
history3=model3.fit(np.array(flatten_train_images),
    ↪to_categorical(train_emotions), epochs=20, batch_size=256,callbacks =
    ↪[time_callback])
time_to_train.append(time_callback.times)
history4=model4.fit(np.array(flatten_train_images),
    ↪to_categorical(train_emotions), epochs=20, batch_size=256,callbacks =
    ↪[time_callback])

```

```

time_to_train.append(time_callback.times)
history5=model5.fit(np.array(flatten_train_images),
    ↳to_categorical(train_emotions), epochs=20, batch_size=256,callbacks =
    ↳[time_callback])
time_to_train.append(time_callback.times)

```

Train image shape: (28709, 2304)  
(28709,)

Epoch 1/20

113/113 [=====] - 4s 11ms/step - loss: 1.6840 -  
accuracy: 0.3310

Epoch 2/20

113/113 [=====] - 1s 11ms/step - loss: 1.5542 -  
accuracy: 0.3973

Epoch 3/20

113/113 [=====] - 1s 10ms/step - loss: 1.4730 -  
accuracy: 0.4306

Epoch 4/20

113/113 [=====] - 1s 10ms/step - loss: 1.4030 -  
accuracy: 0.4601

Epoch 5/20

113/113 [=====] - 1s 10ms/step - loss: 1.3249 -  
accuracy: 0.4910

Epoch 6/20

113/113 [=====] - 1s 10ms/step - loss: 1.2336 -  
accuracy: 0.5309

Epoch 7/20

113/113 [=====] - 1s 10ms/step - loss: 1.1446 -  
accuracy: 0.5661

Epoch 8/20

113/113 [=====] - 1s 10ms/step - loss: 1.0407 -  
accuracy: 0.6076

Epoch 9/20

113/113 [=====] - 1s 10ms/step - loss: 0.9452 -  
accuracy: 0.6470

Epoch 10/20

113/113 [=====] - 1s 10ms/step - loss: 0.8322 -  
accuracy: 0.6881

Epoch 11/20

113/113 [=====] - 1s 10ms/step - loss: 0.7347 -  
accuracy: 0.7276

Epoch 12/20

113/113 [=====] - 1s 10ms/step - loss: 0.6557 -  
accuracy: 0.7582

Epoch 13/20

113/113 [=====] - 1s 10ms/step - loss: 0.5660 -  
accuracy: 0.7953

Epoch 14/20



113/113 [=====] - 1s 10ms/step - loss: 0.4861 -  
 accuracy: 0.8202  
 Epoch 15/20  
 113/113 [=====] - 1s 10ms/step - loss: 0.4244 -  
 accuracy: 0.8468  
 Epoch 16/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.3649 -  
 accuracy: 0.8711  
 Epoch 17/20  
 113/113 [=====] - 1s 10ms/step - loss: 0.3311 -  
 accuracy: 0.8817  
 Epoch 18/20  
 113/113 [=====] - 1s 10ms/step - loss: 0.2946 -  
 accuracy: 0.8944  
 Epoch 19/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.2479 -  
 accuracy: 0.9122  
 Epoch 20/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.2313 -  
 accuracy: 0.9204  
 Epoch 1/20  
 113/113 [=====] - 2s 10ms/step - loss: 1.9349 -  
 accuracy: 0.2956  
 Epoch 2/20  
 113/113 [=====] - 1s 10ms/step - loss: 1.7017 -  
 accuracy: 0.3464  
 Epoch 3/20  
 113/113 [=====] - 1s 10ms/step - loss: 1.6120 -  
 accuracy: 0.3824  
 Epoch 4/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.5344 -  
 accuracy: 0.4082  
 Epoch 5/20  
 113/113 [=====] - 1s 10ms/step - loss: 1.4647 -  
 accuracy: 0.4392  
 Epoch 6/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.4076 -  
 accuracy: 0.4607  
 Epoch 7/20  
 113/113 [=====] - 1s 10ms/step - loss: 1.3538 -  
 accuracy: 0.4807  
 Epoch 8/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.2961 -  
 accuracy: 0.5060  
 Epoch 9/20  
 113/113 [=====] - 1s 10ms/step - loss: 1.2475 -  
 accuracy: 0.5244  
 Epoch 10/20

113/113 [=====] - 1s 10ms/step - loss: 1.1840 -  
 accuracy: 0.5502  
 Epoch 11/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.1234 -  
 accuracy: 0.5747  
 Epoch 12/20  
 113/113 [=====] - 1s 10ms/step - loss: 1.0667 -  
 accuracy: 0.5981  
 Epoch 13/20  
 113/113 [=====] - 1s 10ms/step - loss: 0.9900 -  
 accuracy: 0.6283  
 Epoch 14/20  
 113/113 [=====] - 1s 10ms/step - loss: 0.9311 -  
 accuracy: 0.6522  
 Epoch 15/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.8564 -  
 accuracy: 0.6791  
 Epoch 16/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.8003 -  
 accuracy: 0.7007  
 Epoch 17/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.7307 -  
 accuracy: 0.7295  
 Epoch 18/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.6490 -  
 accuracy: 0.7594  
 Epoch 19/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.5808 -  
 accuracy: 0.7840  
 Epoch 20/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.5216 -  
 accuracy: 0.8090  
 Epoch 1/20  
 113/113 [=====] - 2s 12ms/step - loss: 1.9894 -  
 accuracy: 0.2895  
 Epoch 2/20  
 113/113 [=====] - 1s 12ms/step - loss: 1.7133 -  
 accuracy: 0.3453  
 Epoch 3/20  
 113/113 [=====] - 1s 12ms/step - loss: 1.6252 -  
 accuracy: 0.3750  
 Epoch 4/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.5295 -  
 accuracy: 0.4118  
 Epoch 5/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.4681 -  
 accuracy: 0.4380  
 Epoch 6/20

113/113 [=====] - 1s 12ms/step - loss: 1.3928 -  
 accuracy: 0.4678  
 Epoch 7/20  
 113/113 [=====] - 1s 12ms/step - loss: 1.3372 -  
 accuracy: 0.4922  
 Epoch 8/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.2934 -  
 accuracy: 0.5095  
 Epoch 9/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.2291 -  
 accuracy: 0.5343  
 Epoch 10/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.1772 -  
 accuracy: 0.5555  
 Epoch 11/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.1051 -  
 accuracy: 0.5798  
 Epoch 12/20  
 113/113 [=====] - 1s 11ms/step - loss: 1.0331 -  
 accuracy: 0.6125  
 Epoch 13/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.9570 -  
 accuracy: 0.6419  
 Epoch 14/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.9059 -  
 accuracy: 0.6613  
 Epoch 15/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.8175 -  
 accuracy: 0.6986  
 Epoch 16/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.7428 -  
 accuracy: 0.7217  
 Epoch 17/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.6854 -  
 accuracy: 0.7454  
 Epoch 18/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.6183 -  
 accuracy: 0.7738  
 Epoch 19/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.5543 -  
 accuracy: 0.7951  
 Epoch 20/20  
 113/113 [=====] - 1s 11ms/step - loss: 0.4955 -  
 accuracy: 0.8173  
 Epoch 1/20  
 113/113 [=====] - 3s 14ms/step - loss: 1.7657 -  
 accuracy: 0.3023  
 Epoch 2/20

113/113 [=====] - 1s 13ms/step - loss: 1.6109 -  
accuracy: 0.3669  
Epoch 3/20  
113/113 [=====] - 1s 12ms/step - loss: 1.5468 -  
accuracy: 0.3968  
Epoch 4/20  
113/113 [=====] - 1s 12ms/step - loss: 1.4955 -  
accuracy: 0.4160  
Epoch 5/20  
113/113 [=====] - 1s 12ms/step - loss: 1.4519 -  
accuracy: 0.4345  
Epoch 6/20  
113/113 [=====] - 1s 12ms/step - loss: 1.4028 -  
accuracy: 0.4595  
Epoch 7/20  
113/113 [=====] - 1s 12ms/step - loss: 1.3504 -  
accuracy: 0.4803  
Epoch 8/20  
113/113 [=====] - 1s 12ms/step - loss: 1.2974 -  
accuracy: 0.4972  
Epoch 9/20  
113/113 [=====] - 1s 12ms/step - loss: 1.2347 -  
accuracy: 0.5253  
Epoch 10/20  
113/113 [=====] - 1s 12ms/step - loss: 1.1801 -  
accuracy: 0.5493  
Epoch 11/20  
113/113 [=====] - 1s 12ms/step - loss: 1.1094 -  
accuracy: 0.5798  
Epoch 12/20  
113/113 [=====] - 1s 12ms/step - loss: 1.0416 -  
accuracy: 0.6036  
Epoch 13/20  
113/113 [=====] - 1s 12ms/step - loss: 0.9708 -  
accuracy: 0.6318  
Epoch 14/20  
113/113 [=====] - 1s 12ms/step - loss: 0.8874 -  
accuracy: 0.6667  
Epoch 15/20  
113/113 [=====] - 1s 12ms/step - loss: 0.8051 -  
accuracy: 0.6959  
Epoch 16/20  
113/113 [=====] - 1s 12ms/step - loss: 0.7204 -  
accuracy: 0.7306  
Epoch 17/20  
113/113 [=====] - 1s 12ms/step - loss: 0.6525 -  
accuracy: 0.7576  
Epoch 18/20

113/113 [=====] - 1s 12ms/step - loss: 0.5795 -  
 accuracy: 0.7865  
 Epoch 19/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.5101 -  
 accuracy: 0.8129  
 Epoch 20/20  
 113/113 [=====] - 1s 12ms/step - loss: 0.4550 -  
 accuracy: 0.8345  
 Epoch 1/20  
 113/113 [=====] - 2s 14ms/step - loss: 2.3903 -  
 accuracy: 0.2903  
 Epoch 2/20  
 113/113 [=====] - 2s 14ms/step - loss: 2.0942 -  
 accuracy: 0.3465  
 Epoch 3/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.9958 -  
 accuracy: 0.3712  
 Epoch 4/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.9026 -  
 accuracy: 0.3957  
 Epoch 5/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.8173 -  
 accuracy: 0.4180  
 Epoch 6/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.7519 -  
 accuracy: 0.4341  
 Epoch 7/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.7040 -  
 accuracy: 0.4471  
 Epoch 8/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.6497 -  
 accuracy: 0.4661  
 Epoch 9/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.6144 -  
 accuracy: 0.4777  
 Epoch 10/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.5756 -  
 accuracy: 0.4917  
 Epoch 11/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.5463 -  
 accuracy: 0.5017  
 Epoch 12/20  
 113/113 [=====] - 2s 14ms/step - loss: 1.5239 -  
 accuracy: 0.5153  
 Epoch 13/20  
 113/113 [=====] - 2s 15ms/step - loss: 1.4955 -  
 accuracy: 0.5262  
 Epoch 14/20

```

113/113 [=====] - 2s 14ms/step - loss: 1.4805 -
accuracy: 0.5371
Epoch 15/20
113/113 [=====] - 2s 14ms/step - loss: 1.4442 -
accuracy: 0.5529
Epoch 16/20
113/113 [=====] - 2s 14ms/step - loss: 1.4307 -
accuracy: 0.5584
Epoch 17/20
113/113 [=====] - 2s 14ms/step - loss: 1.4015 -
accuracy: 0.5806
Epoch 18/20
113/113 [=====] - 2s 14ms/step - loss: 1.3664 -
accuracy: 0.5989
Epoch 19/20
113/113 [=====] - 2s 14ms/step - loss: 1.3507 -
accuracy: 0.6087
Epoch 20/20
113/113 [=====] - 2s 14ms/step - loss: 1.3278 -
accuracy: 0.6270

```

**1.2.7 (c.i) (3 points)** Experiment on the validation set with different FNN hyper-parameters, e.g. layers, nodes per layer, activation function, dropout, weight regularization, etc. For each hyper-parameter combination that you have used, please report the following: (1) emotion classification accuracy on the training and validation sets; (2) running time for training the FNN; (3) parameters for each FNN. For 2-3 hyper-parameter combinations, please also plot the cross-entropy loss over the number of iterations during training. Note: If running the FNN takes a long time, you can subsample the input images to a smaller size (e.g., 24 x 24).

```

[11]: # print (time_to_train)
      # for i in time_to_train:
      #     print (sum(i))

validation_emotions=np.array(list(validation_data['emotion']))
validation_pixels=np.array(list(validation_data['pixels_int']))

flatten_validation_images = validation_pixels
validation_performances=[]
models=[model1,model2,model3,model4,model5]
model_names=['model1','model2','model3','model4','model5']
histories=[history1,history2,history3,history4,history5]
for i in range(len(models)):
    print (" ")
    print ("Required details for {}".format(model_names[i]))
    print (" ")

```

```

performance1 = models[i].evaluate(flatten_train_images,
→to_categorical(train_emotions))
print("Emotion Classification Accuracy on the Training set: {0}".
→format(performance1[1]))
print (" ")
performance2 = models[i].evaluate(flatten_validation_images,
→to_categorical(validation_emotions))
validation_performances.append(performance2[1])
print("Emotion Classification Accuracy on the Validation set: {0}".
→format(performance2[1]))
print (" ")
print ("Running time for training the FNN: {} ms".
→format(str(time_to_train[i][0])))
print (" ")
print ("Parameters for the model:")
print (" ")
print (models[i].get_config())
print (" ")
print (models[i].summary())

print (" ")
print('Number of Epochs used to train the model: ', len(histories[i].
→history['loss']))

print(" ")
# print(history.history.keys())
# summarize history for loss
plt.plot(histories[i].history['loss'])
plt.title('{} loss vs epochs'.format(model_names[i]))
plt.ylabel('cross-entropy loss')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(histories[i].history['accuracy'])
plt.title('{} accuracy vs epochs'.format(model_names[i]))
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()

print (" ")
print (" ")

```

```

# performance = model1.evaluate(flatten_test_images,
    ↳to_categorical(validation_emotions))
# print("Accuracy on Test samples: {0}".format(performance[1]))
# performance = model2.evaluate(flatten_test_images,
    ↳to_categorical(validation_emotions))
# print("Accuracy on Test samples: {0}".format(performance[1]))
# performance = model3.evaluate(flatten_test_images,
    ↳to_categorical(validation_emotions))
# print("Accuracy on Test samples: {0}".format(performance[1]))
# performance = model4.evaluate(flatten_test_images,
    ↳to_categorical(validation_emotions))
# print("Accuracy on Test samples: {0}".format(performance[1]))
# performance = model5.evaluate(flatten_test_images,
    ↳to_categorical(validation_emotions))
# print("Accuracy on Test samples: {0}".format(performance[1]))

```

Required details for model1

898/898 [=====] - 4s 4ms/step - loss: 0.1703 -  
accuracy: 0.9428

Emotion Classification Accuracy on the Training set: 0.9428402185440063

113/113 [=====] - 1s 4ms/step - loss: 3.3851 -  
accuracy: 0.4734

Emotion Classification Accuracy on the Validation set: 0.4733909070491791

Running time for training the FNN: 26.956640243530273 ms

Parameters for the model:

```

{'name': 'sequential', 'layers': [{'class_name': 'InputLayer', 'config':
{'batch_input_shape': (None, 2304), 'dtype': 'float32', 'sparse': False,
'ragged': False, 'name': 'first_hidden_layer_input'}}, {'class_name': 'Dense',
'config': {'name': 'first_hidden_layer', 'trainable': True, 'batch_input_shape':
(None, 2304), 'dtype': 'float32', 'units': 1568, 'activation': 'relu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dense', 'config': {'name': 'second_hidden_layer',
'trainable': True, 'dtype': 'float32', 'units': 784, 'activation': 'relu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':

```



```

None}}, {'class_name': 'Dense', 'config': {'name': 'third_hidden_layer',
'trainable': True, 'dtype': 'float32', 'units': 392, 'activation': 'relu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dense', 'config': {'name': 'fourth_hidden_layer',
'trainable': True, 'dtype': 'float32', 'units': 196, 'activation': 'relu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dense', 'config': {'name': 'dense', 'trainable': True,
'dtype': 'float32', 'units': 7, 'activation': 'softmax', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}}]

```

Model: "sequential"

| Layer (type)                | Output Shape | Param # |
|-----------------------------|--------------|---------|
| first_hidden_layer (Dense)  | (None, 1568) | 3614240 |
| second_hidden_layer (Dense) | (None, 784)  | 1230096 |
| third_hidden_layer (Dense)  | (None, 392)  | 307720  |
| fourth_hidden_layer (Dense) | (None, 196)  | 77028   |
| dense (Dense)               | (None, 7)    | 1379    |



```

=====
Total params: 5,230,463
Trainable params: 5,230,463
Non-trainable params: 0

```

-----  
None

Number of Epochs used to train the model: 20



Required details for model2

898/898 [=====] - 4s 5ms/step - loss: 0.3391 -  
accuracy: 0.8836

Emotion Classification Accuracy on the Training set: 0.8835556507110596

113/113 [=====] - 1s 5ms/step - loss: 2.2094 -  
accuracy: 0.4544

Emotion Classification Accuracy on the Validation set: 0.45444414019584656

Running time for training the FNN: 24.502355337142944 ms

Parameters for the model:

```
{'name': 'sequential_1', 'layers': [{'class_name': 'InputLayer', 'config':  
{'batch_input_shape': (None, 2304), 'dtype': 'float32', 'sparse': False,  
'ragged': False, 'name': 'first_hidden_layer_input'}}, {'class_name': 'Dense',  
'config': {'name': 'first_hidden_layer', 'trainable': True, 'batch_input_shape':  
(None, 2304), 'dtype': 'float32', 'units': 1568, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'second_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 784, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'third_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 392, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dropout', 'config': {'name': 'dropout', 'trainable':  
True, 'dtype': 'float32', 'rate': 0.25, 'noise_shape': None, 'seed': None}},  
{'class_name': 'Dense', 'config': {'name': 'dense_1', 'trainable': True,  
'dtype': 'float32', 'units': 7, 'activation': 'softmax', 'use_bias': True,  
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},  
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':  
None, 'bias_regularizer': None, 'activity_regularizer': None,  
'kernel_constraint': None, 'bias_constraint': None}}]}
```

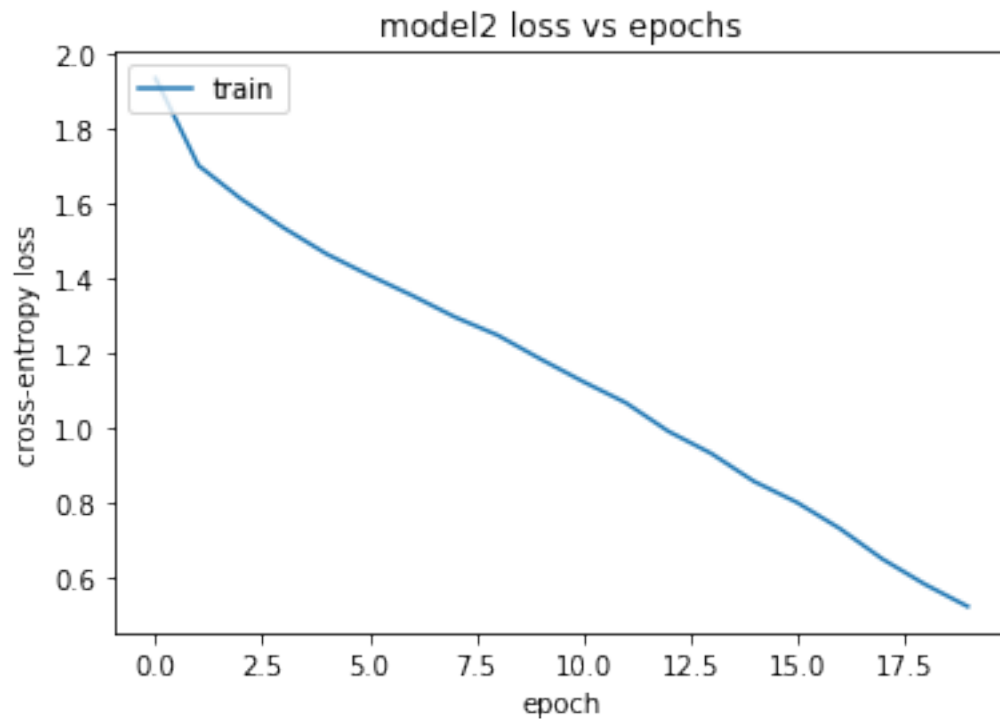
Model: "sequential\_1"


| Layer (type)                | Output Shape | Param # |
|-----------------------------|--------------|---------|
| first_hidden_layer (Dense)  | (None, 1568) | 3614240 |
| second_hidden_layer (Dense) | (None, 784)  | 1230096 |
| third_hidden_layer (Dense)  | (None, 392)  | 307720  |
| dropout (Dropout)           | (None, 392)  | 0       |
| dense_1 (Dense)             | (None, 7)    | 2751    |

Total params: 5,154,807  
 Trainable params: 5,154,807  
 Non-trainable params: 0

None

Number of Epochs used to train the model: 20





Required details for model3

898/898 [=====] - 4s 5ms/step - loss: 0.2900 - accuracy: 0.9070

Emotion Classification Accuracy on the Training set: 0.9069629907608032

113/113 [=====] - 1s 5ms/step - loss: 2.2753 - accuracy: 0.4695

Emotion Classification Accuracy on the Validation set: 0.469490110874176

Running time for training the FNN: 26.87890362739563 ms

Parameters for the model:

```
{'name': 'sequential_2', 'layers': [{'class_name': 'InputLayer', 'config': {'batch_input_shape': (None, 2304), 'dtype': 'float32', 'sparse': False, 'ragged': False, 'name': 'first_hidden_layer_input'}}, {'class_name': 'Dense', 'config': {'name': 'first_hidden_layer', 'trainable': True, 'batch_input_shape': (None, 2304), 'dtype': 'float32', 'units': 2000, 'activation': 'elu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config': {'seed': None}}}]}
```

```
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dense', 'config': {'name': 'second_hidden_layer',
'trainable': True, 'dtype': 'float32', 'units': 1000, 'activation': 'elu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dense', 'config': {'name': 'third_hidden_layer',
'trainable': True, 'dtype': 'float32', 'units': 500, 'activation': 'elu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dropout', 'config': {'name': 'dropout_1', 'trainable':
True, 'dtype': 'float32', 'rate': 0.25, 'noise_shape': None, 'seed': None}},
{'class_name': 'Dense', 'config': {'name': 'dense_2', 'trainable': True,
'dtype': 'float32', 'units': 7, 'activation': 'softmax', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}}]}
```



Model: "sequential\_2"

| Layer (type)                | Output Shape | Param # |
|-----------------------------|--------------|---------|
| first_hidden_layer (Dense)  | (None, 2000) | 4610000 |
| second_hidden_layer (Dense) | (None, 1000) | 2001000 |
| third_hidden_layer (Dense)  | (None, 500)  | 500500  |
| dropout_1 (Dropout)         | (None, 500)  | 0       |
| dense_2 (Dense)             | (None, 7)    | 3507    |

```
=====  
Total params: 7,115,007  
Trainable params: 7,115,007  
Non-trainable params: 0
```

```
-----  
None
```

Number of Epochs used to train the model: 20



Required details for model4

898/898 [=====] - 5s 5ms/step - loss: 0.3394 -  
accuracy: 0.8814

Emotion Classification Accuracy on the Training set: 0.8813612461090088

113/113 [=====] - 1s 5ms/step - loss: 2.5049 -  
accuracy: 0.4531

Emotion Classification Accuracy on the Validation set: 0.45305100083351135

Running time for training the FNN: 29.2670316696167 ms

Parameters for the model:

```
{'name': 'sequential_3', 'layers': [{'class_name': 'InputLayer', 'config':  
{'batch_input_shape': (None, 2304), 'dtype': 'float32', 'sparse': False,  
'ragged': False, 'name': 'first_hidden_layer_input'}}, {'class_name': 'Dense',  
'config': {'name': 'first_hidden_layer', 'trainable': True, 'batch_input_shape':  
(None, 2304), 'dtype': 'float32', 'units': 2000, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'second_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 1000, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'third_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 500, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'fourth_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 250, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'fifth_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 784, 'activation': 'relu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
```


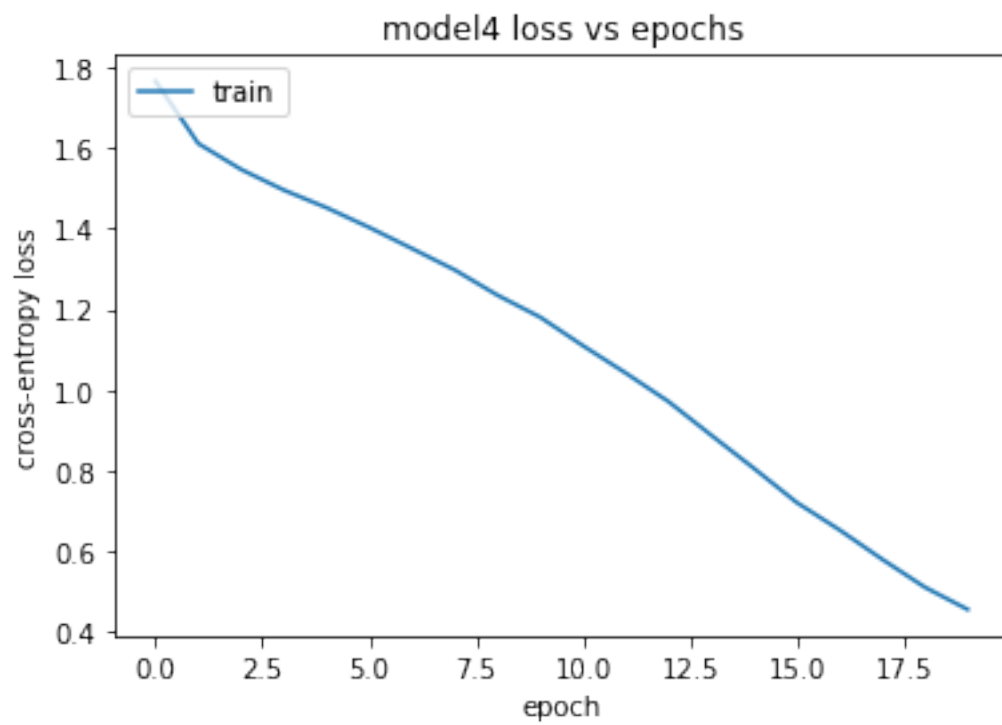


```
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dropout', 'config': {'name': 'dropout_2', 'trainable':
True, 'dtype': 'float32', 'rate': 0.25, 'noise_shape': None, 'seed': None}},
{'class_name': 'Dense', 'config': {'name': 'dense_3', 'trainable': True,
'dtype': 'float32', 'units': 7, 'activation': 'softmax', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}}]}
```

Model: "sequential\_3"

| Layer (type)                | Output Shape | Param # |
|-----------------------------|--------------|---------|
| first_hidden_layer (Dense)  | (None, 2000) | 4610000 |
| second_hidden_layer (Dense) | (None, 1000) | 2001000 |
| third_hidden_layer (Dense)  | (None, 500)  | 500500  |
| fourth_hidden_layer (Dense) | (None, 250)  | 125250  |
| fifth_hidden_layer (Dense)  | (None, 784)  | 196784  |
| dropout_2 (Dropout)         | (None, 784)  | 0       |
| dense_3 (Dense)             | (None, 7)    | 5495    |
| Total params: 7,439,029     |              |         |
| Trainable params: 7,439,029 |              |         |
| Non-trainable params: 0     |              |         |
| None                        |              |         |

Number of Epochs used to train the model: 20



Required details for model5

898/898 [=====] - 5s 6ms/step - loss: 1.2149 -  
accuracy: 0.6733

Emotion Classification Accuracy on the Training set: 0.6732732057571411

113/113 [=====] - 1s 6ms/step - loss: 1.9507 -  
accuracy: 0.4478

Emotion Classification Accuracy on the Validation set: 0.4477570354938507

Running time for training the FNN: 32.77064371109009 ms

Parameters for the model:

```
{'name': 'sequential_4', 'layers': [{'class_name': 'InputLayer', 'config':  
{'batch_input_shape': (None, 2304), 'dtype': 'float32', 'sparse': False,  
'ragged': False, 'name': 'first_hidden_layer_input'}}, {'class_name': 'Dense',  
'config': {'name': 'first_hidden_layer', 'trainable': True, 'batch_input_shape':  
(None, 2304), 'dtype': 'float32', 'units': 2000, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': {'class_name': 'L2', 'config': {'l2':  
9.999999747378752e-05}}, 'bias_regularizer': None, 'activity_regularizer': None,  
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Dense',  
'config': {'name': 'second_hidden_layer', 'trainable': True, 'dtype': 'float32',  
'units': 1000, 'activation': 'elu', 'use_bias': True, 'kernel_initializer':  
{'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer':  
{'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': {'class_name':  
'L2', 'config': {'l2': 9.999999747378752e-05}}, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dense', 'config': {'name': 'third_hidden_layer',  
'trainable': True, 'dtype': 'float32', 'units': 500, 'activation': 'elu',  
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',  
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':  
{}}, 'kernel_regularizer': {'class_name': 'L2', 'config': {'l2':  
9.999999747378752e-05}}, 'bias_regularizer': None, 'activity_regularizer': None,  
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Dense',  
'config': {'name': 'fourth_hidden_layer', 'trainable': True, 'dtype': 'float32',  
'units': 250, 'activation': 'elu', 'use_bias': True, 'kernel_initializer':  
{'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer':  
{'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': {'class_name':  
'L2', 'config': {'l2': 9.999999747378752e-05}}, 'bias_regularizer': None,  
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':  
None}}, {'class_name': 'Dropout', 'config': {'name': 'dropout_3', 'trainable':  
True, 'dtype': 'float32', 'rate': 0.25, 'noise_shape': None, 'seed': None}},
```

```
{'class_name': 'Dense', 'config': {'name': 'dense_4', 'trainable': True,
'dtype': 'float32', 'units': 7, 'activation': 'softmax', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}}]}
```

Model: "sequential\_4"

| Layer (type)                | Output Shape | Param # |
|-----------------------------|--------------|---------|
| first_hidden_layer (Dense)  | (None, 2000) | 4610000 |
| second_hidden_layer (Dense) | (None, 1000) | 2001000 |
| third_hidden_layer (Dense)  | (None, 500)  | 500500  |
| fourth_hidden_layer (Dense) | (None, 250)  | 125250  |
| dropout_3 (Dropout)         | (None, 250)  | 0       |
| dense_4 (Dense)             | (None, 7)    | 1757    |

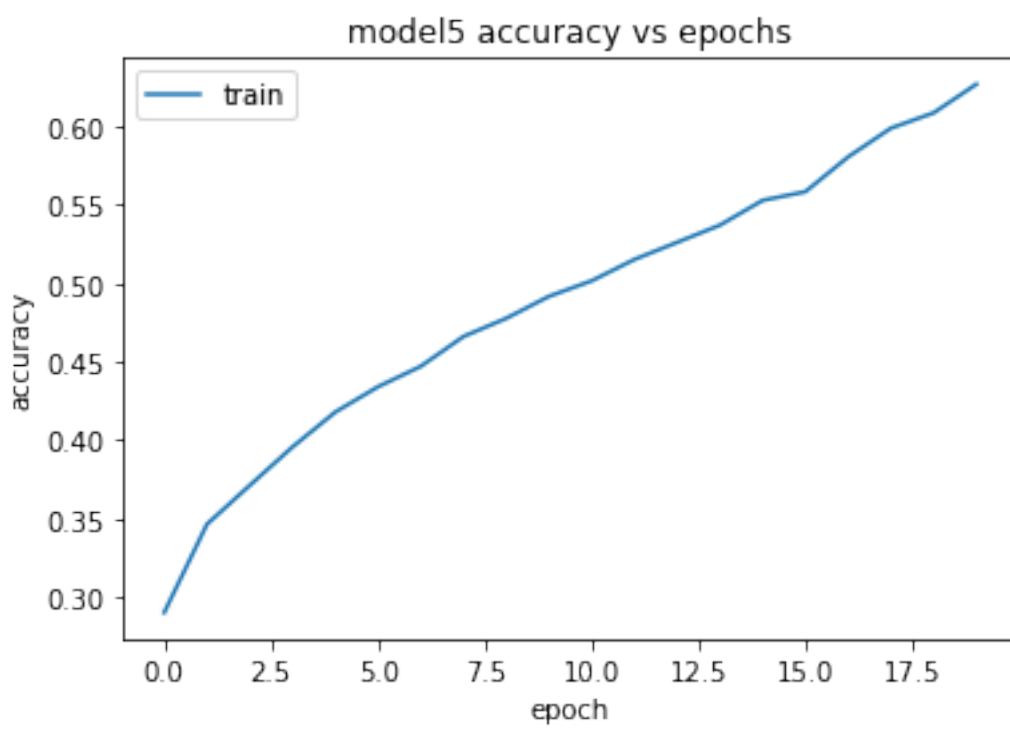

Total params: 7,238,507

Trainable params: 7,238,507

Non-trainable params: 0

None

Number of Epochs used to train the model: 20



### 1.2.8 Answer:

Five different FNN models have been trained on the training dataset and evaluated on the validation dataset. The required details are shown above. The values in a tabular form is shown below:

| Model name | Accuracy on Train dataset (%) | Accuracy on Validation dataset (%) | Training Time (ms) | Parameters Count |
|------------|-------------------------------|------------------------------------|--------------------|------------------|
| model1     | 94.28                         | 47.34                              | 26.95              | 5,230,463        |
| model2     | 88.36                         | 45.44                              | 24.50              | 5,154,807        |
| model3     | 90.70                         | 46.95                              | 26.87              | 7,115,007        |
| model4     | 88.14                         | 45.31                              | 29.26              | 7,439,029        |
| model5     | 67.33                         | 44.78                              | 32.77              | 7,238,507        |

**1.2.9 (c.ii) (1 point) Run the best model that was found based on the validation set from question (c.i) on the testing set. Report the emotion classification accuracy on the testing set.**

### 1.2.10 Answer:

**1.2.11 The best model that was found based on the validation set is:**

```
[12]: # print (validation_performances)
max_validation=validation_performances.index(max(validation_performances))
print (model_names[max_validation])
print (models[max_validation].summary())

print("Emotion Classification Accuracy on the Validation set for the model: {}".format(validation_performances[max_validation]))
```

model1

Model: "sequential"

| Layer (type)                | Output Shape | Param # |
|-----------------------------|--------------|---------|
| first_hidden_layer (Dense)  | (None, 1568) | 3614240 |
| second_hidden_layer (Dense) | (None, 784)  | 1230096 |
| third_hidden_layer (Dense)  | (None, 392)  | 307720  |
| fourth_hidden_layer (Dense) | (None, 196)  | 77028   |
| dense (Dense)               | (None, 7)    | 1379    |

Total params: 5,230,463

Trainable params: 5,230,463

Non-trainable params: 0

-----  
None

Emotion Classification Accuracy on the Validation set for the model:

0.4733909070491791

**1.2.12 The emotion classification accuracy of the model on the testing dataset is:**

```
[13]: test_emotions=np.array(list(test_data['emotion']))
test_pixels=np.array(list(test_data['pixels_int']))

flatten_test_images = test_pixels

test_performance = models[max_validation].evaluate(flatten_test_images,
→to_categorical(test_emotions))
print("Emotion Classification Accuracy on the Testing set: {0}".
→format(test_performance[1]))
```

113/113 [=====] - 1s 5ms/step - loss: 3.2297 -  
accuracy: 0.4728

Emotion Classification Accuracy on the Testing set: 0.47283366322517395

**1.2.13 (d) (4 points) Image classification with CNNs:** In this part, you will use a convolutional neural network (CNN) to perform the emotion classification task.

```
[14]: from tensorflow.keras.layers import Conv2D, Flatten, MaxPooling2D

common_features_1 = [Conv2D(32, kernel_size=3, activation='relu',
→input_shape=(48,48,1)),
    Conv2D(32, kernel_size=3, activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Conv2D(64, kernel_size=3, activation='relu'),
    Conv2D(64, kernel_size=3, activation='relu'),
    MaxPooling2D(pool_size=(2,2)), Flatten(),]
classifier_1 = [Dense(512, activation='relu'), Dense(len(emotions),
→activation='softmax'),]

cnn_model_1 = Sequential(common_features_1+classifier_1)

common_features_2 = [Conv2D(64, kernel_size=3, activation='relu',
→input_shape=(48,48,1),use_bias=True),
    Conv2D(64, kernel_size=3, activation='relu',use_bias=True),
    MaxPooling2D(pool_size=(2,2)),
```

```

        Conv2D(128, kernel_size=3, activation='relu',use_bias=True,
        ↪kernel_regularizer =tf.keras.regularizers.l2( l=0.001)),
        Conv2D(128, kernel_size=3, activation='relu',use_bias=True,
        ↪kernel_regularizer =tf.keras.regularizers.l2( l=0.001)),
        MaxPooling2D(pool_size=(2,2)), Flatten(),]
classifier_2 = [Dense(512, activation='relu',use_bias=True, kernel_regularizer
        ↪=tf.keras.regularizers.l2( l=0.01)), Dense(len(emotions),
        ↪activation='softmax',use_bias=True),]

cnn_model_2 = Sequential(common_features_2+classifier_2)

common_features_3 = [Conv2D(64, kernel_size=3,
        ↪activation='elu',input_shape=(48,48,1),use_bias=True),
        Conv2D(64, kernel_size=3,
        ↪activation='elu',padding='same',use_bias=True),
        MaxPooling2D(pool_size=(2,2)),
        Conv2D(128, kernel_size=3,
        ↪activation='elu',padding='same',use_bias=True),
        MaxPooling2D(pool_size=(2,2)),
        Conv2D(128, kernel_size=3, activation='elu', strides=(2,
        ↪2),padding='same',use_bias=True, kernel_regularizer =tf.keras.regularizers.
        ↪l1( l=0.001)),
        Conv2D(128, kernel_size=3, activation='elu', strides=(2,
        ↪2),padding='same',use_bias=True, kernel_regularizer =tf.keras.regularizers.
        ↪l1( l=0.001)),
        MaxPooling2D(pool_size=(2,2)), Flatten(),]
classifier_3 = [Dense(512, activation='elu',use_bias=True, kernel_regularizer
        ↪=tf.keras.regularizers.l1( l=0.01)),Dropout(0.25), Dense(len(emotions),
        ↪activation='softmax',use_bias=True),]

cnn_model_3 = Sequential(common_features_3+classifier_3)

```

[15]: *# print(cnn\_model\_1.summary()) # Compare number of parameteres against FFN*

```

time_callback = TimeHistory()
cnn_model_1.compile(optimizer='adam',
        ↪loss='categorical_crossentropy',metrics=['accuracy'],)
cnn_model_2.compile(optimizer='adam',
        ↪loss='categorical_crossentropy',metrics=['accuracy'],)
cnn_model_3.compile(optimizer='adam',
        ↪loss='categorical_crossentropy',metrics=['accuracy'],)
time_to_train_cnn=[]

train_images_3d = flatten_train_images.
        ↪reshape(len(flatten_train_images),48,48,1)

```



```

test_images_3d = flatten_test_images.reshape(len(flatten_test_images),48,48,1)

cnn_history_1=cnn_model_1.fit(train_images_3d, to_categorical(train_emotions),
    ↳epochs=15, batch_size=256,callbacks = [time_callback])
time_to_train_cnn.append(time_callback.times)
cnn_history_2=cnn_model_2.fit(train_images_3d, to_categorical(train_emotions),
    ↳epochs=20, batch_size=256,callbacks = [time_callback])
time_to_train_cnn.append(time_callback.times)
cnn_history_3=cnn_model_3.fit(train_images_3d, to_categorical(train_emotions),
    ↳epochs=20, batch_size=256,callbacks = [time_callback])
time_to_train_cnn.append(time_callback.times)

```

Epoch 1/15

6/113 [>...] - ETA: 7s - loss: 1.8666 - accuracy:  
0.2181WARNING:tensorflow:Callback method `on\_train\_batch\_end` is slow compared  
to the batch time (batch time: 0.0308s vs `on\_train\_batch\_end` time: 0.0325s).  
Check your callbacks.

113/113 [=====] - 36s 70ms/step - loss: 1.6138 -  
accuracy: 0.3642

Epoch 2/15

113/113 [=====] - 8s 69ms/step - loss: 1.3458 -  
accuracy: 0.4901

Epoch 3/15

113/113 [=====] - 8s 69ms/step - loss: 1.1907 -  
accuracy: 0.5498

Epoch 4/15

113/113 [=====] - 8s 70ms/step - loss: 1.0795 -  
accuracy: 0.5944

Epoch 5/15

113/113 [=====] - 8s 69ms/step - loss: 0.9614 -  
accuracy: 0.6458

Epoch 6/15

113/113 [=====] - 8s 68ms/step - loss: 0.8132 -  
accuracy: 0.7065

Epoch 7/15

113/113 [=====] - 8s 68ms/step - loss: 0.6542 -  
accuracy: 0.7669

Epoch 8/15

113/113 [=====] - 8s 68ms/step - loss: 0.4602 -  
accuracy: 0.8402

Epoch 9/15

113/113 [=====] - 8s 69ms/step - loss: 0.2866 -  
accuracy: 0.9036

Epoch 10/15

113/113 [=====] - 8s 68ms/step - loss: 0.1756 -  
accuracy: 0.9444

```

Epoch 11/15
113/113 [=====] - 8s 69ms/step - loss: 0.1143 -
accuracy: 0.9666
Epoch 12/15
113/113 [=====] - 8s 68ms/step - loss: 0.0887 -
accuracy: 0.9782
Epoch 13/15
113/113 [=====] - 8s 69ms/step - loss: 0.0805 -
accuracy: 0.9792
Epoch 14/15
113/113 [=====] - 8s 69ms/step - loss: 0.0640 -
accuracy: 0.9859
Epoch 15/15
113/113 [=====] - 8s 69ms/step - loss: 0.0452 -
accuracy: 0.9910
Epoch 1/20
 6/113 [>...] - ETA: 14s - loss: 9.2549 - accuracy:
0.2122WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared
to the batch time (batch time: 0.0470s vs `on_train_batch_end` time: 0.0726s).
Check your callbacks.
113/113 [=====] - 18s 137ms/step - loss: 2.4422 -
accuracy: 0.3194
Epoch 2/20
113/113 [=====] - 15s 135ms/step - loss: 1.6292 -
accuracy: 0.4212
Epoch 3/20
113/113 [=====] - 15s 135ms/step - loss: 1.5323 -
accuracy: 0.4656
Epoch 4/20
113/113 [=====] - 15s 135ms/step - loss: 1.4733 -
accuracy: 0.4891
Epoch 5/20
113/113 [=====] - 15s 135ms/step - loss: 1.4385 -
accuracy: 0.5104
Epoch 6/20
113/113 [=====] - 15s 134ms/step - loss: 1.4073 -
accuracy: 0.5235
Epoch 7/20
113/113 [=====] - 15s 134ms/step - loss: 1.3780 -
accuracy: 0.5388
Epoch 8/20
113/113 [=====] - 15s 134ms/step - loss: 1.3672 -
accuracy: 0.5519
Epoch 9/20
113/113 [=====] - 15s 134ms/step - loss: 1.3364 -
accuracy: 0.5667
Epoch 10/20
113/113 [=====] - 15s 134ms/step - loss: 1.3279 -

```

```

accuracy: 0.5757
Epoch 11/20
113/113 [=====] - 15s 135ms/step - loss: 1.3068 -
accuracy: 0.5918
Epoch 12/20
113/113 [=====] - 15s 135ms/step - loss: 1.2874 -
accuracy: 0.6038
Epoch 13/20
113/113 [=====] - 15s 135ms/step - loss: 1.2769 -
accuracy: 0.6114
Epoch 14/20
113/113 [=====] - 15s 135ms/step - loss: 1.2487 -
accuracy: 0.6278
Epoch 15/20
113/113 [=====] - 15s 135ms/step - loss: 1.2384 -
accuracy: 0.6353
Epoch 16/20
113/113 [=====] - 15s 135ms/step - loss: 1.2262 -
accuracy: 0.6502
Epoch 17/20
113/113 [=====] - 15s 135ms/step - loss: 1.2112 -
accuracy: 0.6625
Epoch 18/20
113/113 [=====] - 15s 134ms/step - loss: 1.2522 -
accuracy: 0.6580
Epoch 19/20
113/113 [=====] - 15s 135ms/step - loss: 1.1952 -
accuracy: 0.6769
Epoch 20/20
113/113 [=====] - 15s 135ms/step - loss: 1.1664 -
accuracy: 0.6935
Epoch 1/20
6/113 [>...] - ETA: 15s - loss: 38.8504 - accuracy:
0.2337WARNING:tensorflow:Callback method `on_train_batch_end` is slow compared
to the batch time (batch time: 0.0521s vs `on_train_batch_end` time: 0.0814s).
Check your callbacks.
113/113 [=====] - 18s 141ms/step - loss: 12.8331 -
accuracy: 0.2555
Epoch 2/20
113/113 [=====] - 16s 138ms/step - loss: 1.9257 -
accuracy: 0.2942
Epoch 3/20
113/113 [=====] - 16s 138ms/step - loss: 1.8848 -
accuracy: 0.3161
Epoch 4/20
113/113 [=====] - 16s 138ms/step - loss: 1.8755 -
accuracy: 0.3279
Epoch 5/20

```

113/113 [=====] - 16s 138ms/step - loss: 1.8652 -  
 accuracy: 0.3353  
 Epoch 6/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.8564 -  
 accuracy: 0.3427  
 Epoch 7/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.8179 -  
 accuracy: 0.3494  
 Epoch 8/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.8348 -  
 accuracy: 0.3554  
 Epoch 9/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.8081 -  
 accuracy: 0.3645  
 Epoch 10/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.7868 -  
 accuracy: 0.3721  
 Epoch 11/20  
 113/113 [=====] - 16s 140ms/step - loss: 1.7774 -  
 accuracy: 0.3831  
 Epoch 12/20  
 113/113 [=====] - 16s 140ms/step - loss: 1.7781 -  
 accuracy: 0.3923  
 Epoch 13/20  
 113/113 [=====] - 16s 140ms/step - loss: 1.7670 -  
 accuracy: 0.3960  
 Epoch 14/20  
 113/113 [=====] - 16s 140ms/step - loss: 1.7509 -  
 accuracy: 0.4016  
 Epoch 15/20  
 113/113 [=====] - 16s 140ms/step - loss: 1.7190 -  
 accuracy: 0.4105  
 Epoch 16/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.7114 -  
 accuracy: 0.4138  
 Epoch 17/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.6906 -  
 accuracy: 0.4200  
 Epoch 18/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.6791 -  
 accuracy: 0.4251  
 Epoch 19/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.6780 -  
 accuracy: 0.4305  
 Epoch 20/20  
 113/113 [=====] - 16s 139ms/step - loss: 1.6748 -  
 accuracy: 0.4347

- 1.2.14 (d.i) (3 points) Experiment on the validation set with different CNN hyper-parameters, e.g. layers, filter size, stride size, activation function, dropout, weight regularization, etc. For each hyper-parameter combination that you have used, please report the following: (1) emotion classification accuracy on the training and validation sets; (2) running time for training the FNN; (3) parameters for each CNN. How do these metrics compare to the FNN?

```
[30]: validation_images_3d = flatten_validation_images.  
      ↪ reshape(len(flatten_validation_images),48,48,1)  
  
validation_performances_cnn=[]  
cnn_models=[cnn_model_1,cnn_model_2,cnn_model_3]  
cnn_model_names=['cnn_model_1','cnn_model_2','cnn_model_3']  
cnn_histories=[cnn_history_1,cnn_history_2,cnn_history_3]  
for i in range(len(cnn_models)):  
    print (" ")  
    print ("Required details for {}".format(cnn_model_names[i]))  
    print (" ")  
    performance1_cnn = cnn_models[i].evaluate(train_images_3d,↵  
    ↪to_categorical(train_emotions))  
    print("Emotion Classification Accuracy on the Training set: {0}".  
    ↪format(performance1_cnn[1]))  
    print (" ")  
    performance2_cnn = cnn_models[i].evaluate(validation_images_3d,↵  
    ↪to_categorical(validation_emotions))  
    validation_performances_cnn.append(performance2_cnn[1])  
    print("Emotion Classification Accuracy on the Validation set: {0}".  
    ↪format(performance2_cnn[1]))  
    print (" ")  
    print ("Running time for training the CNN: {} ms".  
    ↪format(str(time_to_train_cnn[i][0])))  
    print (" ")  
    print ("Parameters for the model:")  
    print (" ")  
    print (cnn_models[i].get_config())  
    print (" ")  
    print (cnn_models[i].summary())  
  
    print (" ")  
    print('Number of Epochs used to train the model: ', len(cnn_histories[i].  
    ↪history['loss']))  
  
    print(" ")  
    # print(history.history.keys())
```

```

# summarize history for loss
plt.plot(cnn_histories[i].history['loss'])
plt.title('{} loss vs epochs'.format(cnn_model_names[i]))
plt.ylabel('cross-entropy loss')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()
# summarize history for accuracy
plt.plot(cnn_histories[i].history['accuracy'])
plt.title('{} accuracy vs epochs'.format(cnn_model_names[i]))
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train'], loc='upper left')
plt.show()

print (" ")
print (" ")

# performance = cnn_model.evaluate(test_images_3d,
    ↪to_categorical(test_emotions))

# print("Accuracy on Test samples: {0}".format(performance[1]))

```

Required details for cnn\_model\_1

898/898 [=====] - 8s 9ms/step - loss: 0.0309 -  
accuracy: 0.9948

Emotion Classification Accuracy on the Training set: 0.9947751760482788

113/113 [=====] - 1s 9ms/step - loss: 2.9449 -  
accuracy: 0.5737

Emotion Classification Accuracy on the Validation set: 0.5736973881721497

Running time for training the CNN: 144.84634375572205 ms

Parameters for the model:

```

{'name': 'sequential_5', 'layers': [{'class_name': 'InputLayer', 'config':
{'batch_input_shape': (None, 48, 48, 1), 'dtype': 'float32', 'sparse': False,
'ragged': False, 'name': 'conv2d_input'}}], {'class_name': 'Conv2D', 'config':
{'name': 'conv2d', 'trainable': True, 'batch_input_shape': (None, 48, 48, 1),
'dtype': 'float32', 'filters': 32, 'kernel_size': (3, 3), 'strides': (1, 1),

```

```

'padding': 'valid', 'data_format': 'channels_last', 'dilation_rate': (1, 1),
'groups': 1, 'activation': 'relu', 'use_bias': True, 'kernel_initializer':
{'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer':
{'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': None,
'bias_regularizer': None, 'activity_regularizer': None, 'kernel_constraint':
None, 'bias_constraint': None}}, {'class_name': 'Conv2D', 'config': {'name':
'conv2d_1', 'trainable': True, 'dtype': 'float32', 'filters': 32, 'kernel_size':
(3, 3), 'strides': (1, 1), 'padding': 'valid', 'data_format': 'channels_last',
'dilation_rate': (1, 1), 'groups': 1, 'activation': 'relu', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name':
'MaxPooling2D', 'config': {'name': 'max_pooling2d', 'trainable': True, 'dtype':
'float32', 'pool_size': (2, 2), 'padding': 'valid', 'strides': (2, 2),
'data_format': 'channels_last'}}, {'class_name': 'Conv2D', 'config': {'name':
'conv2d_2', 'trainable': True, 'dtype': 'float32', 'filters': 64, 'kernel_size':
(3, 3), 'strides': (1, 1), 'padding': 'valid', 'data_format': 'channels_last',
'dilation_rate': (1, 1), 'groups': 1, 'activation': 'relu', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Conv2D',
'config': {'name': 'conv2d_3', 'trainable': True, 'dtype': 'float32', 'filters':
64, 'kernel_size': (3, 3), 'strides': (1, 1), 'padding': 'valid', 'data_format':
'channels_last', 'dilation_rate': (1, 1), 'groups': 1, 'activation': 'relu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'MaxPooling2D', 'config': {'name': 'max_pooling2d_1',
'trainable': True, 'dtype': 'float32', 'pool_size': (2, 2), 'padding': 'valid',
'strides': (2, 2), 'data_format': 'channels_last'}}, {'class_name': 'Flatten',
'config': {'name': 'flatten', 'trainable': True, 'dtype': 'float32',
'data_format': 'channels_last'}}, {'class_name': 'Dense', 'config': {'name':
'dense_5', 'trainable': True, 'dtype': 'float32', 'units': 512, 'activation':
'relu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Dense', 'config': {'name': 'dense_6', 'trainable': True,
'dtype': 'float32', 'units': 7, 'activation': 'softmax', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}}]

```

Model: "sequential\_5"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 46, 46, 32) | 320     |
| conv2d_1 (Conv2D)              | (None, 44, 44, 32) | 9248    |
| max_pooling2d (MaxPooling2D)   | (None, 22, 22, 32) | 0       |
| conv2d_2 (Conv2D)              | (None, 20, 20, 64) | 18496   |
| conv2d_3 (Conv2D)              | (None, 18, 18, 64) | 36928   |
| max_pooling2d_1 (MaxPooling2D) | (None, 9, 9, 64)   | 0       |
| flatten (Flatten)              | (None, 5184)       | 0       |
| dense_5 (Dense)                | (None, 512)        | 2654720 |
| dense_6 (Dense)                | (None, 7)          | 3591    |

Total params: 2,723,303



Trainable params: 2,723,303

Non-trainable params: 0

None

Number of Epochs used to train the model: 15





Required details for cnn\_model\_2

898/898 [=====] - 10s 11ms/step - loss: 1.0117 - accuracy: 0.7653

Emotion Classification Accuracy on the Training set: 0.7652652263641357

113/113 [=====] - 1s 11ms/step - loss: 1.5836 - accuracy: 0.5587

Emotion Classification Accuracy on the Validation set: 0.5586514472961426

Running time for training the CNN: 307.31507754325867 ms

Parameters for the model:

```
{'name': 'sequential_6', 'layers': [{'class_name': 'InputLayer', 'config': {'batch_input_shape': (None, 48, 48, 1), 'dtype': 'float32', 'sparse': False, 'ragged': False, 'name': 'conv2d_4_input'}}, {'class_name': 'Conv2D', 'config': {'name': 'conv2d_4', 'trainable': True, 'batch_input_shape': (None, 48, 48, 1), 'dtype': 'float32', 'filters': 64, 'kernel_size': (3, 3), 'strides': (1, 1), 'padding': 'valid', 'data_format': 'channels_last', 'dilation_rate': (1, 1), 'groups': 1, 'activation': 'relu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': None, 'bias_regularizer': None, 'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Conv2D', 'config': {'name': 'conv2d_5', 'trainable': True, 'dtype': 'float32', 'filters': 64, 'kernel_size': (3, 3), 'strides': (1, 1), 'padding': 'valid', 'data_format': 'channels_last', 'dilation_rate': (1, 1), 'groups': 1, 'activation': 'relu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': None, 'bias_regularizer': None, 'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'MaxPooling2D', 'config': {'name': 'max_pooling2d_2', 'trainable': True, 'dtype': 'float32', 'pool_size': (2, 2), 'padding': 'valid', 'strides': (2, 2), 'data_format': 'channels_last'}}, {'class_name': 'Conv2D', 'config': {'name': 'conv2d_6', 'trainable': True, 'dtype': 'float32', 'filters': 128, 'kernel_size': (3, 3), 'strides': (1, 1), 'padding': 'valid', 'data_format': 'channels_last', 'dilation_rate': (1, 1), 'groups': 1, 'activation': 'relu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': {'class_name': 'L2', 'config': {'l2': 0.0010000000474974513}}, 'bias_regularizer': None, 'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Conv2D', 'config': {'name': 'conv2d_7', 'trainable': True, 'dtype': 'float32', 'filters': 128, 'kernel_size': (3, 3), 'strides': (1, 1), 'padding': 'valid',
```

```

'data_format': 'channels_last', 'dilation_rate': (1, 1), 'groups': 1,
'activation': 'relu', 'use_bias': True, 'kernel_initializer': {'class_name':
'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name':
'Zeros', 'config': {}}, 'kernel_regularizer': {'class_name': 'L2', 'config':
{'l2': 0.0010000000474974513}}, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'MaxPooling2D', 'config': {'name': 'max_pooling2d_3',
'trainable': True, 'dtype': 'float32', 'pool_size': (2, 2), 'padding': 'valid',
'strides': (2, 2), 'data_format': 'channels_last'}}, {'class_name': 'Flatten',
'config': {'name': 'flatten_1', 'trainable': True, 'dtype': 'float32',
'data_format': 'channels_last'}}, {'class_name': 'Dense', 'config': {'name':
'dense_7', 'trainable': True, 'dtype': 'float32', 'units': 512, 'activation':
'relu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': {'class_name': 'L2', 'config': {'l2':
0.009999999776482582}}, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Dense',
'config': {'name': 'dense_8', 'trainable': True, 'dtype': 'float32', 'units': 7,
'activation': 'softmax', 'use_bias': True, 'kernel_initializer': {'class_name':
'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name':
'Zeros', 'config': {}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}}]

```

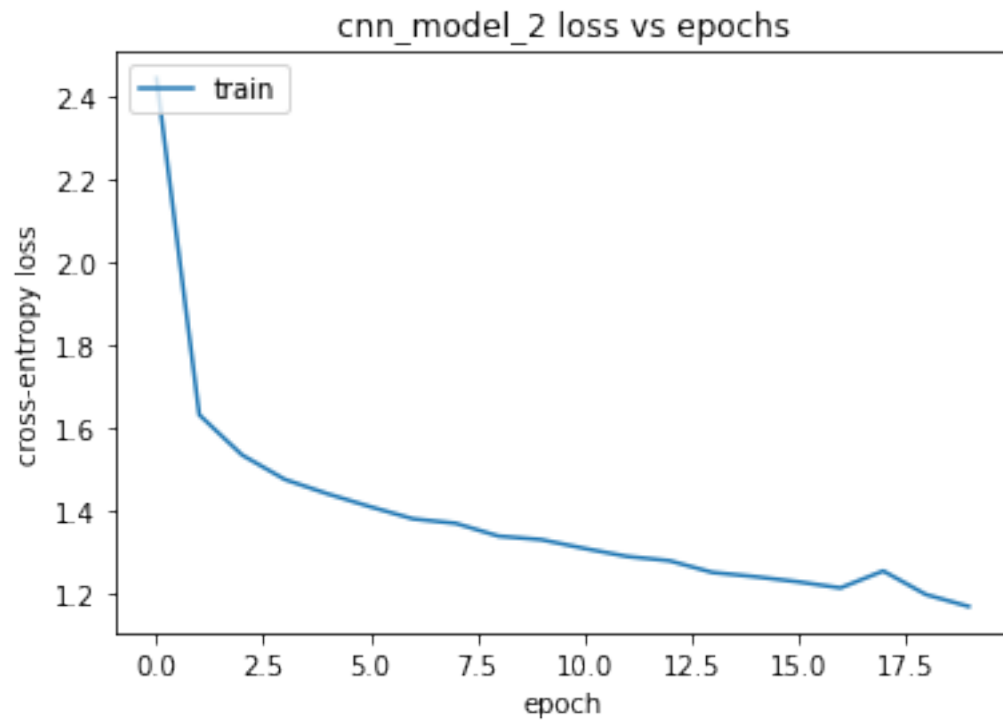
Model: "sequential\_6"


| Layer (type)                    | Output Shape        | Param # |
|---------------------------------|---------------------|---------|
| conv2d_4 (Conv2D)               | (None, 46, 46, 64)  | 640     |
| conv2d_5 (Conv2D)               | (None, 44, 44, 64)  | 36928   |
| max_pooling2d_2 (MaxPooling 2D) | (None, 22, 22, 64)  | 0       |
| conv2d_6 (Conv2D)               | (None, 20, 20, 128) | 73856   |
| conv2d_7 (Conv2D)               | (None, 18, 18, 128) | 147584  |
| max_pooling2d_3 (MaxPooling 2D) | (None, 9, 9, 128)   | 0       |
| flatten_1 (Flatten)             | (None, 10368)       | 0       |
| dense_7 (Dense)                 | (None, 512)         | 5308928 |
| dense_8 (Dense)                 | (None, 7)           | 3591    |

```
=====
Total params: 5,571,527
Trainable params: 5,571,527
Non-trainable params: 0
```

```
-----
None
```

Number of Epochs used to train the model: 20





Required details for cnn\_model\_3

898/898 [=====] - 10s 11ms/step - loss: 1.6395 - accuracy: 0.4426

Emotion Classification Accuracy on the Training set: 0.442613810300827

113/113 [=====] - 1s 11ms/step - loss: 1.6722 - accuracy: 0.4271

Emotion Classification Accuracy on the Validation set: 0.42713847756385803

Running time for training the CNN: 316.84715700149536 ms

Parameters for the model:

```
{'name': 'sequential_7', 'layers': [{'class_name': 'InputLayer', 'config':
{'batch_input_shape': (None, 48, 48, 1), 'dtype': 'float32', 'sparse': False,
'ragged': False, 'name': 'conv2d_8_input'}}, {'class_name': 'Conv2D', 'config':
{'name': 'conv2d_8', 'trainable': True, 'batch_input_shape': (None, 48, 48, 1),
'dtype': 'float32', 'filters': 64, 'kernel_size': (3, 3), 'strides': (1, 1),
'padding': 'valid', 'data_format': 'channels_last', 'dilation_rate': (1, 1),
'groups': 1, 'activation': 'elu', 'use_bias': True, 'kernel_initializer':
```

```

{'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer':
{'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': None,
'bias_regularizer': None, 'activity_regularizer': None, 'kernel_constraint':
None, 'bias_constraint': None}}, {'class_name': 'Conv2D', 'config': {'name':
'conv2d_9', 'trainable': True, 'dtype': 'float32', 'filters': 64, 'kernel_size':
(3, 3), 'strides': (1, 1), 'padding': 'same', 'data_format': 'channels_last',
'dilation_rate': (1, 1), 'groups': 1, 'activation': 'elu', 'use_bias': True,
'kernel_initializer': {'class_name': 'GlorotUniform', 'config': {'seed': None}},
'bias_initializer': {'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer':
None, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name':
'MaxPooling2D', 'config': {'name': 'max_pooling2d_4', 'trainable': True,
'dtype': 'float32', 'pool_size': (2, 2), 'padding': 'valid', 'strides': (2, 2),
'data_format': 'channels_last'}}, {'class_name': 'Conv2D', 'config': {'name':
'conv2d_10', 'trainable': True, 'dtype': 'float32', 'filters': 128,
'kernel_size': (3, 3), 'strides': (1, 1), 'padding': 'same', 'data_format':
'channels_last', 'dilation_rate': (1, 1), 'groups': 1, 'activation': 'elu',
'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':
{}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'MaxPooling2D', 'config': {'name': 'max_pooling2d_5',
'trainable': True, 'dtype': 'float32', 'pool_size': (2, 2), 'padding': 'valid',
'strides': (2, 2), 'data_format': 'channels_last'}}, {'class_name': 'Conv2D',
'config': {'name': 'conv2d_11', 'trainable': True, 'dtype': 'float32',
'filters': 128, 'kernel_size': (3, 3), 'strides': (2, 2), 'padding': 'same',
'data_format': 'channels_last', 'dilation_rate': (1, 1), 'groups': 1,
'activation': 'elu', 'use_bias': True, 'kernel_initializer': {'class_name':
'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name':
'Zeros', 'config': {}}, 'kernel_regularizer': {'class_name': 'L1', 'config':
{'l1': 0.0010000000474974513}}, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'Conv2D', 'config': {'name': 'conv2d_12', 'trainable':
True, 'dtype': 'float32', 'filters': 128, 'kernel_size': (3, 3), 'strides': (2,
2), 'padding': 'same', 'data_format': 'channels_last', 'dilation_rate': (1, 1),
'groups': 1, 'activation': 'elu', 'use_bias': True, 'kernel_initializer':
{'class_name': 'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer':
{'class_name': 'Zeros', 'config': {}}, 'kernel_regularizer': {'class_name':
'L1', 'config': {'l1': 0.00100000000474974513}}, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}, {'class_name': 'MaxPooling2D', 'config': {'name': 'max_pooling2d_6',
'trainable': True, 'dtype': 'float32', 'pool_size': (2, 2), 'padding': 'valid',
'strides': (2, 2), 'data_format': 'channels_last'}}, {'class_name': 'Flatten',
'config': {'name': 'flatten_2', 'trainable': True, 'dtype': 'float32',
'data_format': 'channels_last'}}, {'class_name': 'Dense', 'config': {'name':
'dense_9', 'trainable': True, 'dtype': 'float32', 'units': 512, 'activation':
'elu', 'use_bias': True, 'kernel_initializer': {'class_name': 'GlorotUniform',
'config': {'seed': None}}, 'bias_initializer': {'class_name': 'Zeros', 'config':

```

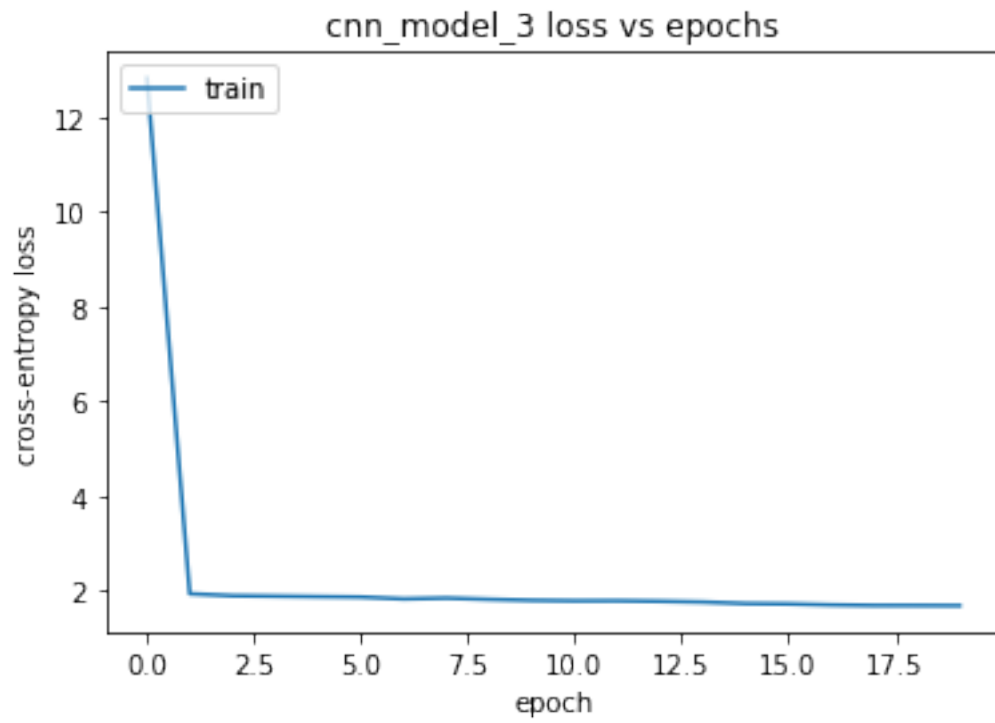
```
{}, 'kernel_regularizer': {'class_name': 'L1', 'config': {'l1':
0.009999999776482582}}, 'bias_regularizer': None, 'activity_regularizer': None,
'kernel_constraint': None, 'bias_constraint': None}}, {'class_name': 'Dropout',
'config': {'name': 'dropout_4', 'trainable': True, 'dtype': 'float32', 'rate':
0.25, 'noise_shape': None, 'seed': None}}, {'class_name': 'Dense', 'config':
{'name': 'dense_10', 'trainable': True, 'dtype': 'float32', 'units': 7,
'activation': 'softmax', 'use_bias': True, 'kernel_initializer': {'class_name':
'GlorotUniform', 'config': {'seed': None}}, 'bias_initializer': {'class_name':
'Zeros', 'config': {}}, 'kernel_regularizer': None, 'bias_regularizer': None,
'activity_regularizer': None, 'kernel_constraint': None, 'bias_constraint':
None}}}]}
```

Model: "sequential\_7"

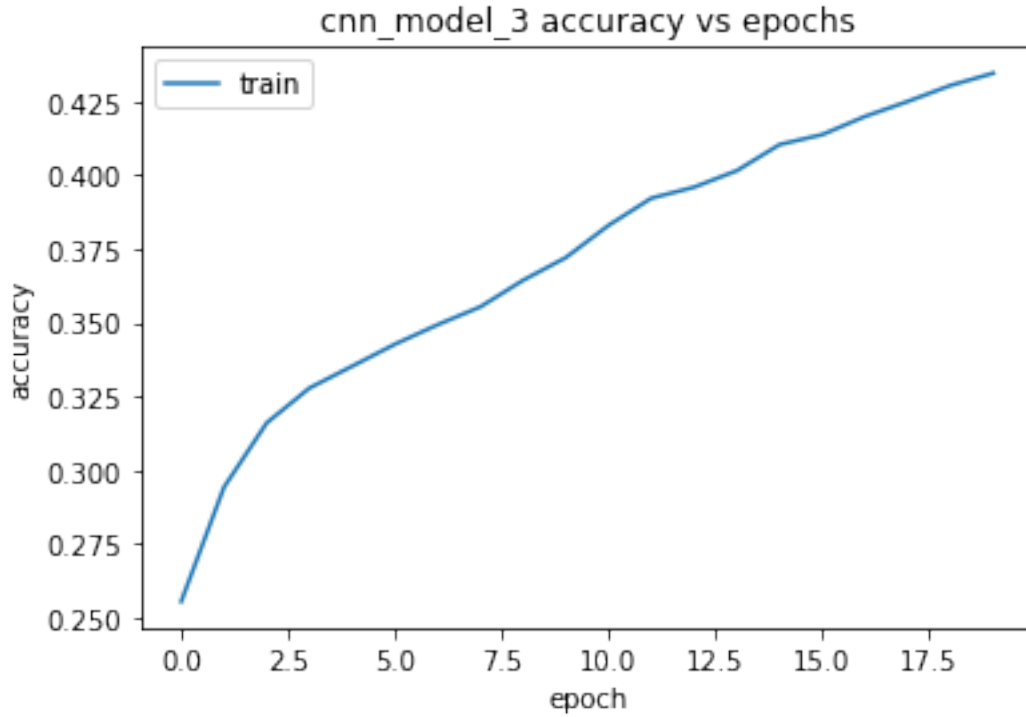
| Layer (type)                    | Output Shape        | Param # |
|---------------------------------|---------------------|---------|
| conv2d_8 (Conv2D)               | (None, 46, 46, 64)  | 640     |
| conv2d_9 (Conv2D)               | (None, 46, 46, 64)  | 36928   |
| max_pooling2d_4 (MaxPooling 2D) | (None, 23, 23, 64)  | 0       |
| conv2d_10 (Conv2D)              | (None, 23, 23, 128) | 73856   |
| max_pooling2d_5 (MaxPooling 2D) | (None, 11, 11, 128) | 0       |
| conv2d_11 (Conv2D)              | (None, 6, 6, 128)   | 147584  |
| conv2d_12 (Conv2D)              | (None, 3, 3, 128)   | 147584  |
| max_pooling2d_6 (MaxPooling 2D) | (None, 1, 1, 128)   | 0       |
| flatten_2 (Flatten)             | (None, 128)         | 0       |
| dense_9 (Dense)                 | (None, 512)         | 66048   |
| dropout_4 (Dropout)             | (None, 512)         | 0       |
| dense_10 (Dense)                | (None, 7)           | 3591    |
| Total params: 476,231           |                     |         |
| Trainable params: 476,231       |                     |         |
| Non-trainable params: 0         |                     |         |

None

Number of Epochs used to train the model: 20







### 1.2.15 Answer:

Three different CNN models have been trained on the training dataset and evaluated on the validation dataset. The required details are shown above. The values in a tabular form is shown below:

| CNN Model number | Accuracy on Train dataset (%) | Accuracy on Validation dataset (%) | Training Time (ms) | Parameters Count |
|------------------|-------------------------------|------------------------------------|--------------------|------------------|
| 1                | 99.48                         | 57.37                              | 144.84 ms          | 2,723,303        |
| 2                | 76.53                         | 55.87                              | 307.31 ms          | 5,571,527        |
| 3                | 44.26                         | 42.71                              | 316.84 ms          | 476,231          |

As can be seen,

1. The accuracy on train, validation, and test dataset is higher in CNN models than in FNN models.
2. The number of parameters to train in a CNN model is also lesser than in FNN model.
3. It is taking more time to train a CNN model than a FNN model.

1.2.16 (d.ii) (1 point) Run the best model that was found based on the validation set from question (d.i) on the testing set. Report the emotion classification accuracy on the testing set. How does this metric compare to the FNN?

1.2.17 Answer:

1.2.18 The best model that was found based on the validation set is:

```
[31]: max_validation_cnn=validation_performances_cnn.  
      ↪index(max(validation_performances_cnn))  
      print (cnn_model_names[max_validation_cnn])  
      print (cnn_models[max_validation_cnn].summary())  
  
      print("Emotion Classification Accuracy on the Validation set for the cnn model:␣  
      ↪{} ".format(validation_performances_cnn[max_validation_cnn]))
```

cnn\_model\_1

Model: "sequential\_5"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 46, 46, 32) | 320     |
| conv2d_1 (Conv2D)              | (None, 44, 44, 32) | 9248    |
| max_pooling2d (MaxPooling2D)   | (None, 22, 22, 32) | 0       |
| conv2d_2 (Conv2D)              | (None, 20, 20, 64) | 18496   |
| conv2d_3 (Conv2D)              | (None, 18, 18, 64) | 36928   |
| max_pooling2d_1 (MaxPooling2D) | (None, 9, 9, 64)   | 0       |
| flatten (Flatten)              | (None, 5184)       | 0       |
| dense_5 (Dense)                | (None, 512)        | 2654720 |
| dense_6 (Dense)                | (None, 7)          | 3591    |

Total params: 2,723,303

Trainable params: 2,723,303

Non-trainable params: 0

None

Emotion Classification Accuracy on the Validation set for the cnn model:

0.5736973881721497

**1.2.19** The emotion classification accuracy of the model on the testing dataset is:

```
[32]: test_images_3d = flatten_test_images.reshape(len(flatten_test_images),48,48,1)

test_performance_cnn = cnn_models[max_validation_cnn].evaluate(test_images_3d,
    ↳to_categorical(test_emotions))
print("Emotion Classification Accuracy on the Testing set for the CNN: {0}".
    ↳format(test_performance_cnn[1]))
print (" ")
print("And Emotion Classification Accuracy on the Testing set for the FNN: {0}".
    ↳format(test_performance[1]))
```

```
113/113 [=====] - 1s 8ms/step - loss: 2.7430 -
accuracy: 0.5812
Emotion Classification Accuracy on the Testing set for the CNN:
0.5812203884124756
```

```
And Emotion Classification Accuracy on the Testing set for the FNN:
0.47283366322517395
```

**1.2.20** We can see that the Emotion Classification Accuracy on the Testing set for the CNN (58.12%) is better than that of the FNN (47.28%).

**1.2.21** (e) (1 point) Bayesian optimization for hyper-parameter tuning: Instead of performing grid or random search to tune the hyper-parameters of the CNN, we can also try a model-based method for finding the optimal hyper-parameters through Bayesian optimization. This method performs a more intelligent search on the hyper-parameter space in order to estimate the best set of hyper-parameters for the data. Use publicly available libraries (e.g., hyperopt in Python) to perform a Bayesian optimization on the hyper-parameter space using the validation set. Report the emotion classification accuracy on the testing set.

```
[19]: !pip install hyperopt
```

```
Requirement already satisfied: hyperopt in /usr/local/lib/python3.7/dist-
packages (0.1.2)
Requirement already satisfied: tqdm in /usr/local/lib/python3.7/dist-packages
(from hyperopt) (4.62.3)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages
(from hyperopt) (1.15.0)
Requirement already satisfied: networkx in /usr/local/lib/python3.7/dist-
packages (from hyperopt) (2.6.3)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages
(from hyperopt) (1.19.5)
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages
(from hyperopt) (1.4.1)
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages
(from hyperopt) (0.16.0)
```

Requirement already satisfied: pymongo in /usr/local/lib/python3.7/dist-packages (from hyperopt) (3.12.1)

```
[20]: from hyperopt import hp, fmin, tpe, STATUS_OK, Trials
```

```
[21]: def optimize_cnn(hyperparameter):

    # Defining model using hyperparameters
    cnn_model = Sequential([Conv2D(32,
    ↪kernel_size=hyperparameter['conv_kernel_size'],
    ↪activation=hyperparameter['activation_fn'], input_shape=(48,48,1)),
        Conv2D(32, kernel_size=hyperparameter['conv_kernel_size'],
    ↪activation=hyperparameter['activation_fn']),
        MaxPooling2D(pool_size=(2,2)),
    ↪Dropout(hyperparameter['dropout_prob']),
        Conv2D(64, kernel_size=hyperparameter['conv_kernel_size'],
    ↪activation=hyperparameter['activation_fn']),
        Conv2D(128, kernel_size=hyperparameter['conv_kernel_size'],
    ↪activation=hyperparameter['activation_fn']),
        MaxPooling2D(pool_size=(2,2)),
    ↪Dropout(hyperparameter['dropout_prob']),
        Flatten(),
        Dense(512, hyperparameter['activation_fn']),
        Dense(len(emotions), activation='softmax'),])

    cnn_model.compile(optimizer=hyperparameter['optimizer'],
    ↪loss='categorical_crossentropy', metrics=['accuracy'],)

    train_X, train_y = train_images_3d, train_emotions
    valid_X, valid_y = validation_images_3d, validation_emotions

    _ = cnn_model.fit(train_X, to_categorical(train_y), epochs=10,
    ↪batch_size=256, verbose=0)
    # Evaluating accuracy on validation data
    performance = cnn_model.evaluate(valid_X, to_categorical(valid_y), verbose=0)

    print("Hyperparameters: ", hyperparameter, "Accuracy: ", performance[1])
    print("-----")

    return({"status": STATUS_OK, "loss": -1*performance[1], "model":cnn_model})

# Defining search space for hyper-parameters
space = {
    # The kernel_size for convolutions:
    'conv_kernel_size': hp.choice('conv_kernel_size', [1, 3, 5]),
```

```

# Uniform distribution in finding appropriate dropout values
'dropout_prob': hp.uniform('dropout_prob', 0.1, 0.35),
# Choice of optimizer
'optimizer': hp.choice('optimizer', ['Adam', 'sgd']),
#choice of activation function
'activation_fn': hp.choice('activation', ['relu', 'sigmoid', 'elu']),
}

trials = Trials()

# Finding the best hyperparameters
best = fmin(
    optimize_cnn,
    space,
    algo=tpe.suggest,
    trials=trials,
    max_evals=15,
)

print("=====")
print("Best Hyperparameters", best)

test_model = trials.results[np.argmin([r['loss'] for r in trials.
    ↳results])]['model']

performance = test_model.evaluate(test_images_3d, to_categorical(test_emotions))

print("=====")
print("Test Accuracy: ", performance[1])

```

Hyperparameters:  
{'activation\_fn': 'sigmoid', 'conv\_kernel\_size': 5, 'dropout\_prob':  
0.15266477923626062, 'optimizer': 'sgd'}

Accuracy:  
0.24937307834625244

-----  
Hyperparameters:  
{'activation\_fn': 'elu', 'conv\_kernel\_size': 3, 'dropout\_prob':  
0.31596639352914324, 'optimizer': 'Adam'}

Accuracy:  
0.5539147257804871

-----  
Hyperparameters:  
{'activation\_fn': 'relu', 'conv\_kernel\_size': 1, 'dropout\_prob':  
0.1744496091134044, 'optimizer': 'Adam'}

Accuracy:  
0.47088325023651123

```

-----
Hyperparameters:
{'activation_fn': 'sigmoid', 'conv_kernel_size': 1, 'dropout_prob':
0.287162391630506, 'optimizer': 'sgd'}
Accuracy:
0.24937307834625244
-----
Hyperparameters:
{'activation_fn': 'relu', 'conv_kernel_size': 5, 'dropout_prob':
0.2448089365657019, 'optimizer': 'Adam'}
Accuracy:
0.592644214630127
-----
Hyperparameters:
{'activation_fn': 'sigmoid', 'conv_kernel_size': 1, 'dropout_prob':
0.29375151885611395, 'optimizer': 'Adam'}
Accuracy:
0.24937307834625244
-----
Hyperparameters:
{'activation_fn': 'elu', 'conv_kernel_size': 3, 'dropout_prob':
0.12498432433670034, 'optimizer': 'sgd'}
Accuracy:
0.3736416697502136
-----
Hyperparameters:
{'activation_fn': 'elu', 'conv_kernel_size': 5, 'dropout_prob':
0.316955992072036, 'optimizer': 'sgd'}
Accuracy:
0.37977153062820435
-----
Hyperparameters:
{'activation_fn': 'sigmoid', 'conv_kernel_size': 3, 'dropout_prob':
0.1439456589291522, 'optimizer': 'sgd'}
Accuracy:
0.1819448322057724
-----
Hyperparameters:
{'activation_fn': 'sigmoid', 'conv_kernel_size': 1, 'dropout_prob':
0.14360078356291128, 'optimizer': 'Adam'}
Accuracy:
0.24937307834625244
-----
Hyperparameters:
{'activation_fn': 'elu', 'conv_kernel_size': 1, 'dropout_prob':
0.14857276563606964, 'optimizer': 'sgd'}
Accuracy:
0.38395094871520996

```

```
-----
Hyperparameters:
{'activation_fn': 'elu', 'conv_kernel_size': 1, 'dropout_prob':
0.21999128824946373, 'optimizer': 'Adam'}
Accuracy:
0.37698522210121155
-----
```

```
Hyperparameters:
{'activation_fn': 'relu', 'conv_kernel_size': 5, 'dropout_prob':
0.27824443302614543, 'optimizer': 'Adam'}
Accuracy:
0.585121214389801
-----
```

```
Hyperparameters:
{'activation_fn': 'relu', 'conv_kernel_size': 1, 'dropout_prob':
0.1268864685218872, 'optimizer': 'Adam'}
Accuracy:
0.4505433142185211
-----
```

```
Hyperparameters:
{'activation_fn': 'elu', 'conv_kernel_size': 3, 'dropout_prob':
0.13194258629013292, 'optimizer': 'Adam'}
Accuracy:
0.5215937495231628
-----
```

```
100%|          | 15/15 [21:24<00:00, 85.63s/it, best loss: -0.592644214630127]
=====
```

```
Best Hyperparameters {'activation': 0, 'conv_kernel_size': 2, 'dropout_prob':
0.2448089365657019, 'optimizer': 0}
113/113 [=====] - 1s 7ms/step - loss: 1.0938 -
accuracy: 0.5940
=====
```

```
Test Accuracy: 0.5940373539924622
```

**1.2.22** The best hypermeters are: {'activation': 'relu', 'conv\_kernel\_size': 5, 'dropout\_prob': 0.2448089365657019, 'optimizer': 'Adam'}

**1.2.23** The emotion classification accuracy on the testing set with the best hyperparameters is as follows:

```
[22]: performance = test_model.evaluate(test_images_3d, to_categorical(test_emotions))

print("=====")
print("Test Accuracy: ", performance[1])
```

```
113/113 [=====] - 1s 8ms/step - loss: 1.0938 -
accuracy: 0.5940
=====
Test Accuracy: 0.5940373539924622
```

- 1.2.24 (f) (Bonus - 1 point) Fine-tuning: Use a pre-trained CNN (e.g., the pre-trained example of the MNIST dataset that we saw in class, or any other available pre-trained CNN) and fine-tune it on the FER data. Please experiment with different fine-tuning hyper-parameters (e.g., layers to fine-tune, regularization during fine-tuning) on the validation set. Report the classification accuracy for all hyper-parameter combinations on the validation set. Also report the classification accuracy with the best hyper-parameter combination on the testing set.

```
[23]: from tensorflow.keras.applications import MobileNetV2
from tensorflow.keras.applications.inception_v3 import InceptionV3
from tensorflow.keras.applications.efficientnet import EfficientNetB0
from tensorflow.keras.applications.resnet50 import ResNet50

from tensorflow.keras.applications.inception_resnet_v2 import InceptionResNetV2
from tensorflow.python.keras.layers import Dense, Flatten,
    ↳GlobalAveragePooling2D, Activation, Flatten, Dropout, BatchNormalization
from keras.applications.inception_resnet_v2 import preprocess_input
from tensorflow.keras.layers import RandomFlip, RandomRotation
from tensorflow.keras import Input, Model
from tensorflow.image import grayscale_to_rgb
from tensorflow import convert_to_tensor

model_ft_1 = Sequential()
model_ft_1.add(ResNet50(input_shape=(48,48,3), include_top=False,
    ↳pooling='avg', weights="imagenet"))
# model.trainable=False
model_ft_1.add(Dense(512))
model_ft_1.add(Activation('relu'))
model_ft_1.add(Dense(1024))
model_ft_1.add(Activation('relu'))
model_ft_1.add(Dense(512))
model_ft_1.add(Activation('relu'))
model_ft_1.add(Dropout(0.3))
model_ft_1.add(Dense(len(emotions), activation='softmax'))
model_ft_1.layers[0].trainable = False
model_ft_1.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

model_ft_2 = Sequential()
model_ft_2.add(MobileNetV2(input_shape=(48,48,3), include_top=False,
    ↳pooling='avg', weights="imagenet"))
# model.trainable=False
model_ft_2.add(Dense(512))
```



```

model_ft_2.add(Activation('relu'))
model_ft_2.add(Dense(1024))
model_ft_2.add(Activation('relu'))
model_ft_2.add(Dense(512))
model_ft_2.add(Activation('relu'))
model_ft_2.add(Dropout(0.4))
model_ft_2.add(Dense(len(emotions), activation='softmax'))
model_ft_2.layers[0].trainable = False
model_ft_2.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

```

```

model_ft_3 = Sequential()
model_ft_3.add(MobileNetV2(input_shape=(48,48,3), include_top=False,
    ↳pooling='avg', weights="imagenet"))
# model.trainable=False
model_ft_3.add(Dense(512, kernel_regularizer=regularizers.l2(0.01)))
model_ft_3.add(Activation('elu'))
model_ft_3.add(Dense(1024, kernel_regularizer=regularizers.l2(0.01)))
model_ft_3.add(Activation('elu'))
model_ft_3.add(Dense(512, kernel_regularizer=regularizers.l2(0.01)))
model_ft_3.add(Activation('elu'))
model_ft_3.add(Dropout(0.3))
model_ft_3.add(Dense(len(emotions), activation='softmax'))
model_ft_3.layers[0].trainable = False
model_ft_3.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

```

```

train_images_3d_3=grayscale_to_rgb(convert_to_tensor(train_images_3d))

history_ft_1 = model_ft_1.fit(train_images_3d_3,
    ↳to_categorical(train_emotions), epochs=10, batch_size=256, verbose=1)
history_ft_2 = model_ft_2.fit(train_images_3d_3,
    ↳to_categorical(train_emotions), epochs=15, batch_size=256, verbose=1)
history_ft_3 = model_ft_3.fit(train_images_3d_3,
    ↳to_categorical(train_emotions), epochs=15, batch_size=256, verbose=1)

```

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_notop.h5](https://storage.googleapis.com/tensorflow/keras-applications/resnet/resnet50_weights_tf_dim_ordering_tf_kernels_notop.h5)  
 94773248/94765736 [=====] - 1s 0us/step  
 94781440/94765736 [=====] - 1s 0us/step  
 WARNING:tensorflow:`input\_shape` is undefined or non-square, or `rows` is not in

[96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

Downloading data from [https://storage.googleapis.com/tensorflow/keras-applications/mobilenet\\_v2/mobilenet\\_v2\\_weights\\_tf\\_dim\\_ordering\\_tf\\_kernels\\_1.0\\_224\\_no\\_top.h5](https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5)

9412608/9406464 [=====] - 0s 0us/step

9420800/9406464 [=====] - 0s 0us/step

WARNING:tensorflow:`input\_shape` is undefined or non-square, or `rows` is not in [96, 128, 160, 192, 224]. Weights for input shape (224, 224) will be loaded as the default.

Epoch 1/10

113/113 [=====] - 19s 123ms/step - loss: 1.8139 - accuracy: 0.2447

Epoch 2/10

113/113 [=====] - 13s 119ms/step - loss: 1.7682 - accuracy: 0.2650

Epoch 3/10

113/113 [=====] - 13s 119ms/step - loss: 1.7483 - accuracy: 0.2810

Epoch 4/10

113/113 [=====] - 13s 119ms/step - loss: 1.7168 - accuracy: 0.3006

Epoch 5/10

113/113 [=====] - 13s 119ms/step - loss: 1.7078 - accuracy: 0.3065

Epoch 6/10

113/113 [=====] - 14s 120ms/step - loss: 1.6895 - accuracy: 0.3203

Epoch 7/10

113/113 [=====] - 14s 120ms/step - loss: 1.6880 - accuracy: 0.3198

Epoch 8/10

113/113 [=====] - 13s 119ms/step - loss: 1.6691 - accuracy: 0.3344

Epoch 9/10

113/113 [=====] - 14s 120ms/step - loss: 1.6706 - accuracy: 0.3336

Epoch 10/10

113/113 [=====] - 14s 120ms/step - loss: 1.6562 - accuracy: 0.3427

Epoch 1/15

113/113 [=====] - 12s 72ms/step - loss: 1.6875 - accuracy: 0.3315

Epoch 2/15

113/113 [=====] - 8s 70ms/step - loss: 1.5157 - accuracy: 0.4143

Epoch 3/15

113/113 [=====] - 8s 69ms/step - loss: 1.3848 -

```

accuracy: 0.4745
Epoch 4/15
113/113 [=====] - 8s 69ms/step - loss: 1.1953 -
accuracy: 0.5571
Epoch 5/15
113/113 [=====] - 8s 70ms/step - loss: 0.9751 -
accuracy: 0.6465
Epoch 6/15
113/113 [=====] - 8s 70ms/step - loss: 0.7550 -
accuracy: 0.7327
Epoch 7/15
113/113 [=====] - 8s 69ms/step - loss: 0.5858 -
accuracy: 0.7967
Epoch 8/15
113/113 [=====] - 8s 69ms/step - loss: 0.4517 -
accuracy: 0.8442
Epoch 9/15
113/113 [=====] - 8s 69ms/step - loss: 0.3559 -
accuracy: 0.8786
Epoch 10/15
113/113 [=====] - 8s 69ms/step - loss: 0.2589 -
accuracy: 0.9131
Epoch 11/15
113/113 [=====] - 8s 69ms/step - loss: 0.2060 -
accuracy: 0.9319
Epoch 12/15
113/113 [=====] - 8s 70ms/step - loss: 0.1797 -
accuracy: 0.9402
Epoch 13/15
113/113 [=====] - 8s 69ms/step - loss: 0.1545 -
accuracy: 0.9500
Epoch 14/15
113/113 [=====] - 8s 69ms/step - loss: 0.1486 -
accuracy: 0.9513
Epoch 15/15
113/113 [=====] - 8s 70ms/step - loss: 0.1208 -
accuracy: 0.9588
Epoch 1/15
113/113 [=====] - 12s 72ms/step - loss: 1.8576 -
accuracy: 0.3234
Epoch 2/15
113/113 [=====] - 8s 72ms/step - loss: 1.5784 -
accuracy: 0.3922
Epoch 3/15
113/113 [=====] - 8s 72ms/step - loss: 1.4853 -
accuracy: 0.4257
Epoch 4/15
113/113 [=====] - 8s 72ms/step - loss: 1.4071 -

```

```

accuracy: 0.4622
Epoch 5/15
113/113 [=====] - 8s 72ms/step - loss: 1.3213 -
accuracy: 0.4960
Epoch 6/15
113/113 [=====] - 8s 72ms/step - loss: 1.2378 -
accuracy: 0.5301
Epoch 7/15
113/113 [=====] - 8s 72ms/step - loss: 1.1497 -
accuracy: 0.5668
Epoch 8/15
113/113 [=====] - 8s 73ms/step - loss: 1.0500 -
accuracy: 0.6091
Epoch 9/15
113/113 [=====] - 8s 72ms/step - loss: 0.9576 -
accuracy: 0.6406
Epoch 10/15
113/113 [=====] - 8s 72ms/step - loss: 0.8568 -
accuracy: 0.6835
Epoch 11/15
113/113 [=====] - 8s 73ms/step - loss: 0.7484 -
accuracy: 0.7248
Epoch 12/15
113/113 [=====] - 8s 72ms/step - loss: 0.6589 -
accuracy: 0.7587
Epoch 13/15
113/113 [=====] - 8s 72ms/step - loss: 0.5760 -
accuracy: 0.7881
Epoch 14/15
113/113 [=====] - 8s 72ms/step - loss: 0.4831 -
accuracy: 0.8257
Epoch 15/15
113/113 [=====] - 8s 71ms/step - loss: 0.4291 -
accuracy: 0.8449

```

```

[24]: validation_images_3d_3=grayscale_to_rgb(convert_to_tensor(validation_images_3d))

# performance = model_ft_3.evaluate(validation_images_3d_3,
#   ↳to_categorical(validation_emotions))

# print("=====")
# print("Test Accuracy: ", performance[1])

validation_performances_ft=[]
ft_models=[model_ft_1,model_ft_2,model_ft_3]
ft_model_names=['model_ft_1','model_ft_2','model_ft_3']

```

```

ft_histories=[history_ft_1,history_ft_2,history_ft_3]
for i in range(len(ft_models)):
    print (" ")
    print ("Required details for {}".format(ft_model_names[i]))
    print (" ")
    performance1_ft = ft_models[i].evaluate(train_images_3d_3,
    ↳to_categorical(train_emotions))
    print("Emotion Classification Accuracy on the Training set: {0}".
    ↳format(performance1_ft[1]))
    print (" ")
    performance2_ft = ft_models[i].evaluate(validation_images_3d_3,
    ↳to_categorical(validation_emotions))
    validation_performances_ft.append(performance2_ft[1])
    print("Emotion Classification Accuracy on the Validation set: {0}".
    ↳format(performance2_ft[1]))
    print (" ")
    print ("Parameters for the model:")
    print (" ")
    print (ft_models[i].summary())

    print (" ")
    print('Number of Epochs used to train the model: ', len(ft_histories[i].
    ↳history['loss']))

    print(" ")
    # print(history.history.keys())
    # summarize history for loss
    plt.plot(ft_histories[i].history['loss'])
    plt.title('{} loss vs epochs'.format(ft_model_names[i]))
    plt.ylabel('cross-entropy loss')
    plt.xlabel('epoch')
    plt.legend(['train'], loc='upper left')
    plt.show()
    # summarize history for accuracy
    plt.plot(ft_histories[i].history['accuracy'])
    plt.title('{} accuracy vs epochs'.format(ft_model_names[i]))
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['train'], loc='upper left')
    plt.show()

    print (" ")
    print (" ")

```

Required details for model\_ft\_1

898/898 [=====] - 33s 35ms/step - loss: 1.6846 - accuracy: 0.3244

Emotion Classification Accuracy on the Training set: 0.32442787289619446

113/113 [=====] - 4s 35ms/step - loss: 1.7151 - accuracy: 0.3104

Emotion Classification Accuracy on the Validation set: 0.3103928565979004

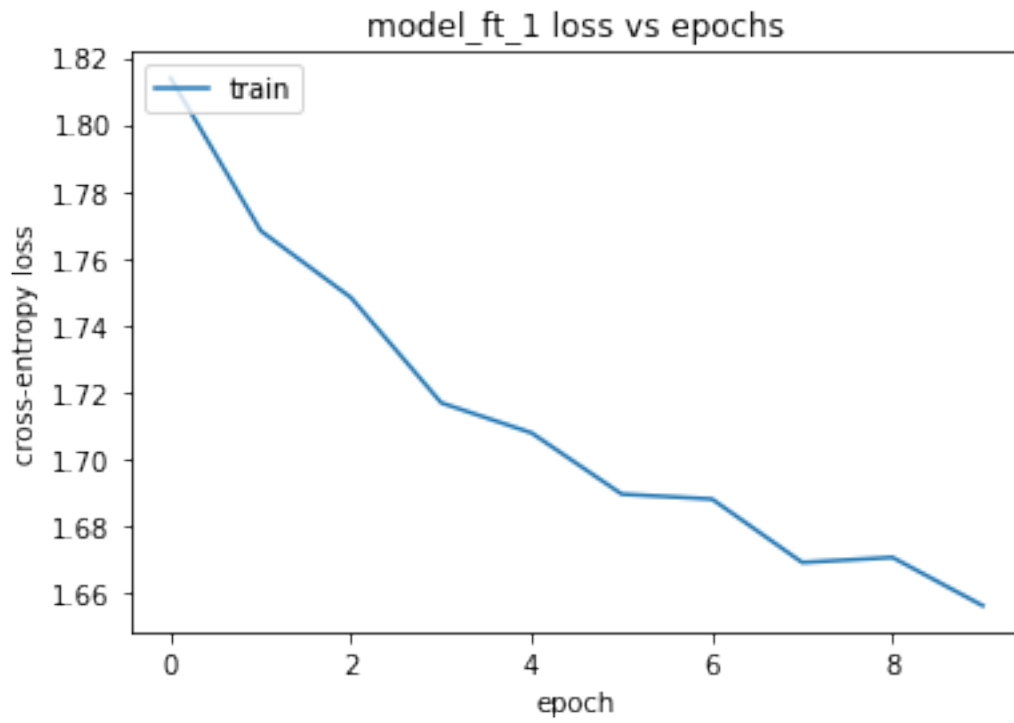
Parameters for the model:


Model: "sequential\_23"

| Layer (type)                     | Output Shape | Param #  |
|----------------------------------|--------------|----------|
| resnet50 (Functional)            | (None, 2048) | 23587712 |
| module_wrapper (ModuleWrapper)   | (None, 512)  | 1049088  |
| module_wrapper_1 (ModuleWrapper) | (None, 512)  | 0        |
| module_wrapper_2 (ModuleWrapper) | (None, 1024) | 525312   |
| module_wrapper_3 (ModuleWrapper) | (None, 1024) | 0        |
| module_wrapper_4 (ModuleWrapper) | (None, 512)  | 524800   |
| module_wrapper_5 (ModuleWrapper) | (None, 512)  | 0        |
| module_wrapper_6 (ModuleWrapper) | (None, 512)  | 0        |
| module_wrapper_7 (ModuleWrapper) | (None, 7)    | 3591     |
| Total params: 25,690,503         |              |          |
| Trainable params: 2,102,791      |              |          |
| Non-trainable params: 23,587,712 |              |          |

None

Number of Epochs used to train the model: 10





Required details for model\_ft\_2

898/898 [=====] - 20s 21ms/step - loss: 0.0624 - accuracy: 0.9808

Emotion Classification Accuracy on the Training set: 0.9807725548744202

113/113 [=====] - 2s 21ms/step - loss: 3.8308 - accuracy: 0.4230

Emotion Classification Accuracy on the Validation set: 0.42295902967453003

Parameters for the model:

Model: "sequential\_24"



| Layer (type)                      | Output Shape | Param # |
|-----------------------------------|--------------|---------|
| mobilenetv2_1.00_224 (Functional) | (None, 1280) | 2257984 |
| module_wrapper_8 (ModuleWrapper)  | (None, 512)  | 655872  |



|                                       |              |        |
|---------------------------------------|--------------|--------|
| module_wrapper_9 (ModuleWr<br>apper)  | (None, 512)  | 0      |
| module_wrapper_10 (ModuleWr<br>apper) | (None, 1024) | 525312 |
| module_wrapper_11 (ModuleWr<br>apper) | (None, 1024) | 0      |
| module_wrapper_12 (ModuleWr<br>apper) | (None, 512)  | 524800 |
| module_wrapper_13 (ModuleWr<br>apper) | (None, 512)  | 0      |
| module_wrapper_14 (ModuleWr<br>apper) | (None, 512)  | 0      |
| module_wrapper_15 (ModuleWr<br>apper) | (None, 7)    | 3591   |

```
=====
Total params: 3,967,559
Trainable params: 1,709,575
Non-trainable params: 2,257,984
-----
None
```

Number of Epochs used to train the model: 15



Required details for model\_ft\_3

898/898 [=====] - 19s 20ms/step - loss: 0.3500 - accuracy: 0.8713

Emotion Classification Accuracy on the Training set: 0.8713295459747314

113/113 [=====] - 2s 19ms/step - loss: 2.8958 - accuracy: 0.4322

Emotion Classification Accuracy on the Validation set: 0.4321537911891937

Parameters for the model:

Model: "sequential\_25"


| Layer (type)                      | Output Shape | Param # |
|-----------------------------------|--------------|---------|
| mobilenetv2_1.00_224 (Functional) | (None, 1280) | 2257984 |
| module_wrapper_16 (ModuleWrapper) | (None, 512)  | 655872  |
| module_wrapper_17 (ModuleWrapper) | (None, 512)  | 0       |
| module_wrapper_18 (ModuleWrapper) | (None, 1024) | 525312  |
| module_wrapper_19 (ModuleWrapper) | (None, 1024) | 0       |
| module_wrapper_20 (ModuleWrapper) | (None, 512)  | 524800  |
| module_wrapper_21 (ModuleWrapper) | (None, 512)  | 0       |
| module_wrapper_22 (ModuleWrapper) | (None, 512)  | 0       |
| module_wrapper_23 (ModuleWrapper) | (None, 7)    | 3591    |


=====  
Total params: 3,967,559

Trainable params: 1,709,575  
Non-trainable params: 2,257,984

-----  
None

Number of Epochs used to train the model: 15





### 1.2.25 Answer:

Three different fine-tuned models have been trained on the training dataset and evaluated on the validation dataset. The required details are shown above. The values in a tabular form is shown below:

| Model name | Accuracy on Train dataset (%) | Accuracy on Validation dataset (%) | Total Trainable Parameters Count |
|------------|-------------------------------|------------------------------------|----------------------------------|
| model_ft_1 | 32.44                         | 31.04                              | 2,102,791                        |
| model_ft_2 | 98.08                         | 42.30                              | 1,709,575                        |
| model_ft_3 | 87.13                         | 43.22                              | 1,709,575                        |

### 1.2.26 The best fine-tuned model that was found based on the validation set is:

```
[42]: max_validation_ft=validation_performances_ft.
      ↪ index(max(validation_performances_ft))
      print (ft_model_names[max_validation_ft])
      print (ft_models[max_validation_ft].summary())
```

```
print("Emotion Classification Accuracy on the Validation set for the fine-tuned_
↪model: {}".format(validation_performances_ft[max_validation_ft]))
```

model\_ft\_3

Model: "sequential\_25"

| Layer (type)                      | Output Shape | Param # |
|-----------------------------------|--------------|---------|
| mobilenetv2_1.00_224 (Functional) | (None, 1280) | 2257984 |
| module_wrapper_16 (ModuleWrapper) | (None, 512)  | 655872  |
| module_wrapper_17 (ModuleWrapper) | (None, 512)  | 0       |
| module_wrapper_18 (ModuleWrapper) | (None, 1024) | 525312  |
| module_wrapper_19 (ModuleWrapper) | (None, 1024) | 0       |
| module_wrapper_20 (ModuleWrapper) | (None, 512)  | 524800  |
| module_wrapper_21 (ModuleWrapper) | (None, 512)  | 0       |
| module_wrapper_22 (ModuleWrapper) | (None, 512)  | 0       |
| module_wrapper_23 (ModuleWrapper) | (None, 7)    | 3591    |

Total params: 3,967,559

Trainable params: 1,709,575

Non-trainable params: 2,257,984

None

Emotion Classification Accuracy on the Validation set for the fine-tuned model:  
0.4321537911891937

**1.2.27** The emotion classification accuracy of the best fine-tuned model (created using mobilenetv2) on the testing dataset is:

```
[43]: test_images_3d_3=grayscale_to_rgb(convert_to_tensor(test_images_3d))

test_performance_ft= ft_models[max_validation_ft].evaluate(test_images_3d_3,
↳to_categorical(test_emotions))
print("Emotion Classification Accuracy on the Testing set for the fine-tuned_
↳model: {0}".format(test_performance_ft[1]))
```

```
113/113 [=====] - 2s 20ms/step - loss: 2.8454 -
accuracy: 0.4235
Emotion Classification Accuracy on the Testing set for the fine-tuned model:
0.42351630330085754
```

**1.2.28** (g) (Bonus - 1 point) Data augmentation: Data augmentation is a way to increase the size of our dataset and reduce overfitting, especially when we use complicated models with many parameters to learn. Using any available toolbox or your own code, implement some of these techniques and augment the original FER data.

```
[41]: from keras.preprocessing.image import ImageDataGenerator

for i in range(9):
    plt.subplot(330 + 1 + i)
    plt.title('Actual image')
    plt.imshow(train_images_3d[i].reshape(48, 48), cmap=plt.get_cmap('gray'))
plt.show()

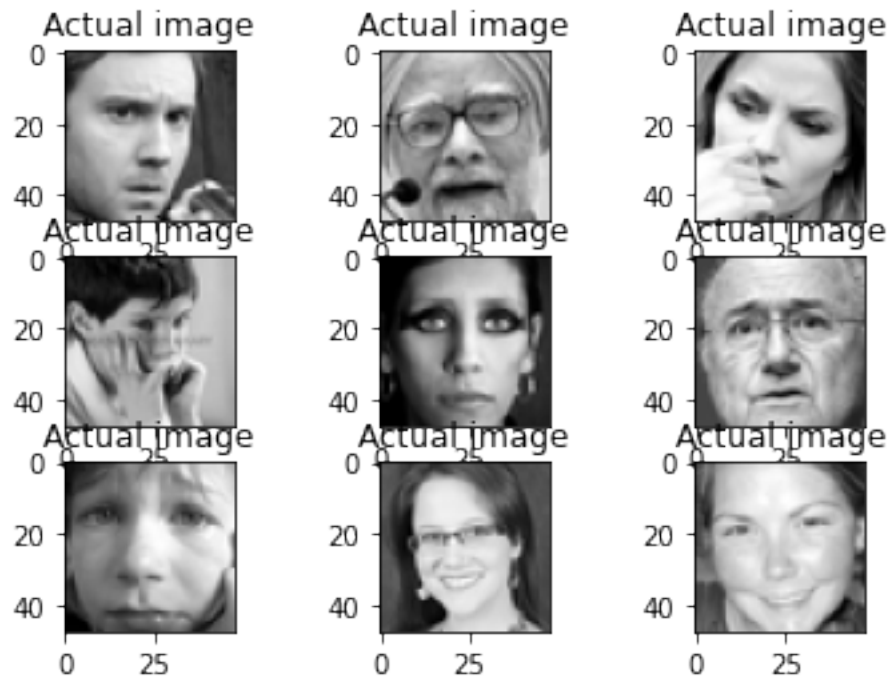
shift=0.2
datagen1 = ImageDataGenerator(featurewise_center=True,
↳featurewise_std_normalization=True, zca_whitening=True, rotation_range=90,
↳width_shift_range=0.2, height_shift_range=0.2, horizontal_flip=True,
↳vertical_flip=True)
datagen2 = ImageDataGenerator(featurewise_center=True,
↳featurewise_std_normalization=True)
datagen3 = ImageDataGenerator(zca_whitening=True)
datagen4 = ImageDataGenerator(rotation_range=90)
datagen5 = ImageDataGenerator(width_shift_range=shift, height_shift_range=shift)

datagens=[datagen1, datagen2, datagen3, datagen4, datagen5]
labels=["Augmentation 1", "Augmentation 2", "Augmentation 3", "Augmentation 4",
↳"Augmentation 5"]
for ij, datagen in enumerate(datagens):
    datagen.fit(train_images_3d[:9])
```

```

for X_batch, y_batch in datagen.flow(train_images_3d[:9], train_emotions[:
→9], batch_size=9):
    for i in range(0, 9):
        plt.subplot(330 + 1 + i)
        plt.title(labels[ij])
        plt.imshow(X_batch[i].reshape(48, 48), cmap=plt.get_cmap('gray'))
    plt.show()
    break

```

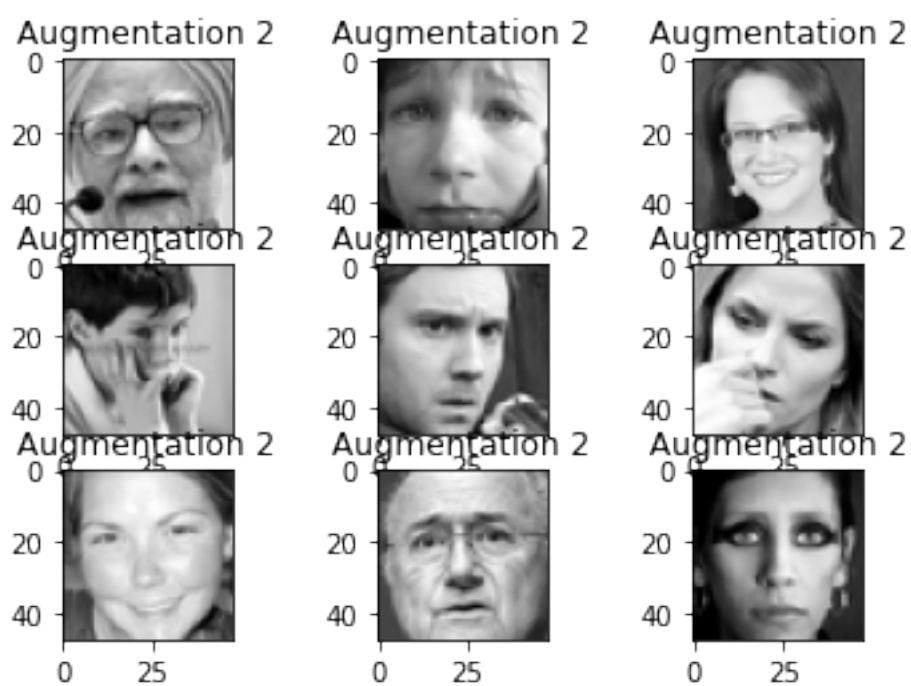



```


/usr/local/lib/python3.7/dist-
packages/keras_preprocessing/image/image_data_generator.py:342: UserWarning:
This ImageDataGenerator specifies `zca_whitening` which overrides setting
of `featurewise_std_normalization`.
    warnings.warn('This ImageDataGenerator specifies '
/usr/local/lib/python3.7/dist-
packages/keras_preprocessing/image/image_data_generator.py:337: UserWarning:
This ImageDataGenerator specifies `zca_whitening`, which overrides setting of
`featurewise_center`.
    warnings.warn('This ImageDataGenerator specifies '

```









**Answer:** Some outputs of the following data augmentation techniques are shown above:

1. datagen1 = ImageDataGenerator(featurewise\_center=True, featurewise\_std\_normalization=True, zca\_whitening=True, rotation\_range=90, width\_shift\_range=0.2, height\_shift\_range=0.2, horizontal\_flip=True, vertical\_flip=True)
2. datagen2 = ImageDataGenerator(featurewise\_center=True, featurewise\_std\_normalization=True)
3. datagen3 = ImageDataGenerator(zca\_whitening=True)
4. datagen4 = ImageDataGenerator(rotation\_range=90)
5. datagen5 = ImageDataGenerator(width\_shift\_range=shift, height\_shift\_range=shift)

Now, training a CNN model using datagen1:

```
[44]: #training a model on augmented dataset
datagen1.fit(train_images_3d)
it = datagen1.flow(train_images_3d, to_categorical(train_emotions))

common_features = [Conv2D(32, kernel_size=3, activation='relu',
    ↪input_shape=(48,48,1)),
    Conv2D(32, kernel_size=3, activation='relu'),
    MaxPooling2D(pool_size=(2,2)),
    Conv2D(64, kernel_size=3, activation='relu'),
    Conv2D(64, kernel_size=3, activation='relu'),
```

```

        MaxPooling2D(pool_size=(2,2)), Flatten(),]
classifier = [Dense(512, activation='relu'), Dense(len(emotions),
↳activation='softmax'),]

cnn_model = Sequential(common_features+classifier)
cnn_model.compile(optimizer='adam',
↳loss='categorical_crossentropy',metrics=['accuracy'],)
history_cnn = cnn_model.fit_generator(it, epochs=10)

```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:16: UserWarning:  
`Model.fit\_generator` is deprecated and will be removed in a future version.  
Please use `Model.fit`, which supports generators.

```
app.launch_new_instance()
```

Epoch 1/10

```
898/898 [=====] - 69s 76ms/step - loss: 1.8030 -
accuracy: 0.2509
```

Epoch 2/10

```
898/898 [=====] - 68s 76ms/step - loss: 1.7720 -
accuracy: 0.2677
```

Epoch 3/10

```
898/898 [=====] - 69s 77ms/step - loss: 1.7104 -
accuracy: 0.3034
```

Epoch 4/10

```
898/898 [=====] - 69s 77ms/step - loss: 1.6771 -
accuracy: 0.3227
```

Epoch 5/10

```
898/898 [=====] - 69s 77ms/step - loss: 1.6565 -
accuracy: 0.3326
```

Epoch 6/10

```
898/898 [=====] - 71s 79ms/step - loss: 1.6429 -
accuracy: 0.3403
```

Epoch 7/10

```
898/898 [=====] - 70s 78ms/step - loss: 1.6280 -
accuracy: 0.3488
```

Epoch 8/10

```
898/898 [=====] - 70s 78ms/step - loss: 1.6102 -
accuracy: 0.3577
```

Epoch 9/10

```
898/898 [=====] - 70s 78ms/step - loss: 1.5926 -
accuracy: 0.3715
```

Epoch 10/10

```
898/898 [=====] - 71s 79ms/step - loss: 1.5837 -
accuracy: 0.3755
```

The accuracy on the test dataset using this model is as follows:

```
[46]: performance = cnn_model.evaluate(test_images_3d, to_categorical(test_emotions))  
  
print("Accuracy on Test samples: {0}".format(performance[1]))
```

```
113/113 [=====] - 1s 9ms/step - loss: 2.1034 -  
accuracy: 0.1677  
Accuracy on Test samples: 0.16773474216461182
```

```
[29]:
```